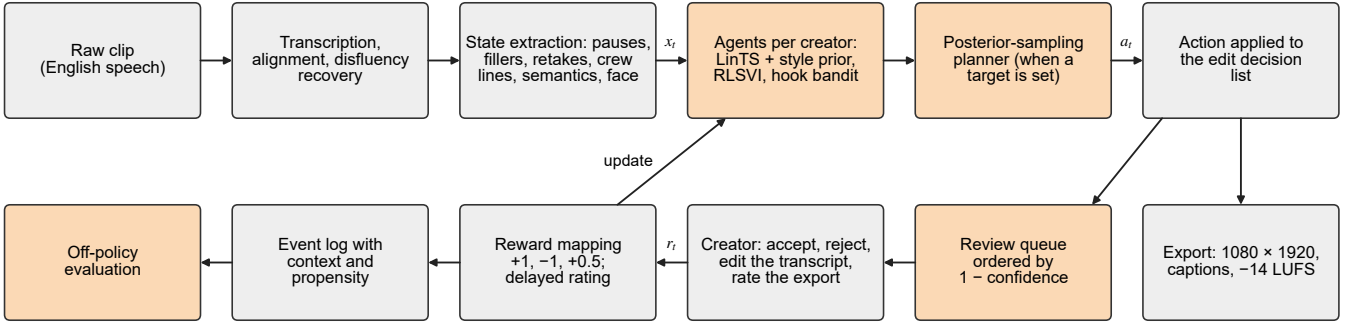
Raw clip
(English speech)
Transcription,
alignment, disfluency
recovery
State extraction: pauses,
fillers, retakes, crew
lines, semantics, face
$x_t$
Agents per creator:
LinTS + style prior,
RLSVI, hook bandit
Posterior-sampling
planner (when a
target is set)
$a_t$
Action applied to
the edit decision
list
update
Off-policy
evaluation
Event log with
context and
propensity
Reward mapping
+1, −1, +0.5;
delayed rating
$r_t$
Creator: accept, reject,
edit the transcript,
rate the export
Review queue
ordered by
1 − confidence
Export: 1080 × 1920,
captions, −14 LUFS