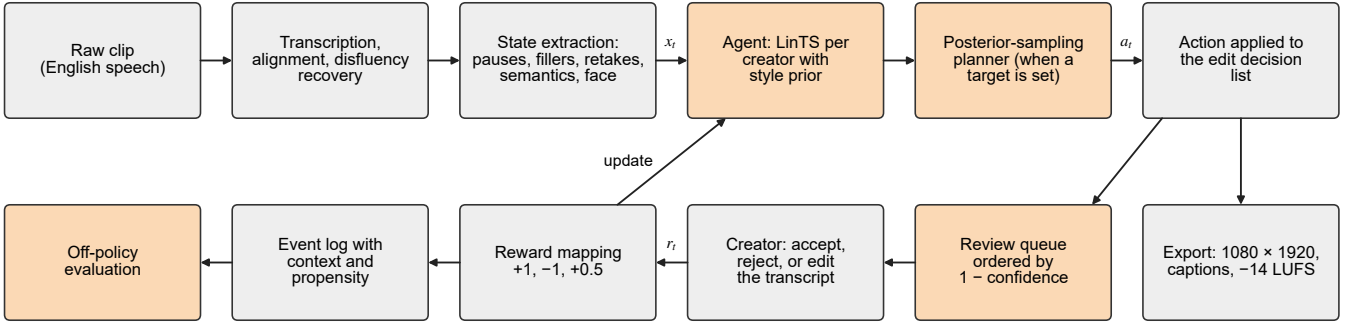
Raw clip
(English speech)
Transcription,
alignment, disfluency
recovery
State extraction:
pauses, fillers, retakes,
semantics, face
x_t
Agent: LinTS per
creator with
style prior
Posterior-sampling
planner (when a
target is set)
a_t
Action applied to
the edit decision
list
update
Off-policy
evaluation
Event log with
context and
propensity
Reward mapping
+1, −1, +0.5
r_t
Creator: accept,
reject, or edit
the transcript
Review queue
ordered by
1 − confidence
Export: 1080 × 1920,
captions, −14 LUFS