IEEE Access
Multidisciplinary | Rapid Review | Open Access Journal

Date of publication xxxx 00, 0000, date of current version xxxx 00, 0000.

Digital Object Identifier 10.1109/ACCESS.2017.DOI

Learning a Creator's Editing Style Online: Posterior Sampling With Style Priors, Uncertainty-Ordered Review, and Length Planning for Talking-Head Video

P. SUPREETH REDDY[1], B. VIJAYASIMHA REDDY[1], M. RUSHIVARUN REDDY[1], and GEETHA M[1]

[1]Department of Computer Science and Engineering, Amrita School of Computing, Amrita Vishwa Vidyapeetham, Amritapuri, Kerala, India (e-mail: am.sc.u4ae23142@am.students.amrita.edu, am.sc.u4aie23110@am.students.amrita.edu, am.sc.u4aie23128@am.students.amrita.edu, geetham@am.amrita.edu)

Corresponding author: P. Supreeth Reddy (e-mail: am.sc.u4ae23142@am.students.amrita.edu).

ABSTRACT Short talking-head videos are cut by hand, and two creators rarely want the same cut from the same footage. Fixed rules cannot capture that difference, and supervised models need a labeled history that a new creator lacks. We treat each edit as a decision in online reinforcement learning, with the creator's accepts, rejections, and ordinary transcript edits as reward. A linear Thompson sampling agent handles the per-phrase decisions. Around it we add a tempered posterior over editing styles for the cold start, a review queue ordered by posterior uncertainty, an exact posterior-sampling planner for target lengths, and a randomized value-iteration agent for visual pacing. Across six simulated creators and ten seeds, the method reached 93.1 percent acceptance by the tenth video, against 58.6 percent for a rule baseline, with 9 percent fewer wrong proposals than the best fairly tuned neural agent. Networks draw level on acceptance by the fiftieth video. The planner cut held-out length error from 40.0 to 4.6 percent for 0.4 points of acceptance. Value iteration trailed a bandit on beat accuracy but earned better ratings when per-beat feedback was scarce. Three recorded clips exercised the analysis pipeline; no formal human study has been run.

INDEX TERMS Contextual bandits, human-in-the-loop learning, off-policy evaluation, personalization, posterior sampling, reinforcement learning, Thompson sampling, video editing

REINFORCEMENT learning (RL) studies agents that improve by acting and observing the consequences of what they did [1]. Unlike supervised learning, the agent gets no list of right answers. It sees a reward for the action it took and has to decide when to try something else in order to learn more. That framing has moved well beyond games [2]: recommender systems [3]–[5], photo retouching [6], and the alignment of large language models with human preference [7], [8] all now use it. In the last two cases the reward comes from people, which makes every sample expensive. Sample efficiency stops being a detail and becomes the main design constraint.

Video editing sits squarely in that expensive-feedback corner. A creator who films a talking-head Reel or Short usually records it in one sitting, restarts sentences, pauses to think, and says "um" and "you know". Getting from that footage to something postable means removing dead air, choosing between repeated takes, deciding which fillers matter, adding punch-in zooms, and highlighting a caption word now and then. Each choice is small, but a two-minute clip contains over a hundred of them. They also depend on taste. A storyteller keeps the pause that lands a joke, whereas a creator who wants a fast pace removes it. So there is no single correct edit against which a universal model could be trained.

Treating this as RL raises several specific difficulties. Feedback is partial, because a rejected pause suggestion says only that this option was wrong, not which of the other three was right. A creator will only edit a handful of videos before forming an opinion of the tool, so the agent has to be useful from the first video. Tastes change over time. Editing errors also differ in cost, since a badly placed cut deletes a sentence

but a misplaced zoom is merely ugly. Finally, a length limit such as "make this 30 seconds" applies to the edit as a whole and cannot be decided one cut at a time.

The task addressed here is the following. Given a raw clip of English speech and a history of one creator's corrections, propose a value for each editing decision in the clip so that the proposals match what that creator would have chosen, ask the creator only about the ones most likely to be wrong, respect an optional target length, and update the model from whatever the creator does next.

Existing tools leave a clear gap. Rule-based features such as one-click silence removal apply one threshold to everyone and never learn. Transcript-based editors [9]–[11] make cutting far faster but leave every decision to the person. Supervised models could learn taste, yet they need labeled data from that same person before they can be of use, and they do not choose what to ask about. RL fits better because the labels arrive only as a side effect of using the tool, and an uncertainty-aware agent can improve while it works. We use the simplest formulation that captures each part. Most editing decisions barely influence later contexts, so a contextual bandit [3], [12] is a defensible model that learns from a few hundred interactions. The sequential parts, the length budget and the pacing of visual effects, get an exact planner and a randomized value-iteration agent. We tested the claim that deep RL [2], [13] needs more data than this setting offers instead of assuming it.

Our technical objectives were to build a simulation environment that exercises the real analysis code, to design a per-creator agent and a reward function that covers both suggestions and manual edits, to outperform rule-based and generic baselines within a few videos, to keep learning stable across seeds, and to measure honestly where the method fails. This paper makes the following contributions.

- **A reward model in which ordinary editing is the training signal.** Word deletions, pause drags, and caption clicks are classified into overrides of learnable decisions, so a creator never performs a labeling step. In the ablation study, removing this mapping lowers early acceptance by 2.5 points.
- **A style posterior as a tempered bandit prior.** A Bayesian mixture over style prototypes supplies the prior mean of the bandit. Tempering lets the displayed style sharpen within about four interactions, whereas the prior moves slowly enough that one click cannot rewrite the policy. With leave-one-style-out priors, regret falls by 7 percent and the median number of decisions needed to reach 85 percent acceptance falls from 430 to 344.
- **Uncertainty-ordered review under a budget.** Thompson draws give each suggestion a confidence (expected calibration error 0.023), and that value sorts the review queue. At a 30 percent review budget, 86.9 percent of exported decisions are correct, against 78.4 percent for random order and 63.3 percent for timeline order.
- **Posterior sampling for a target-length edit with guarded reward.** We pose length fitting as a finite-horizon Markov decision process (MDP) whose reward is the learned bandit model, solve it exactly, and show that the history features must be frozen to stop the solver from exploiting model extrapolation. Length error on unseen videos falls from 40.0 to 4.6 percent.
- **A measured account of what exploration and sequential models add here.** Policies trained with and without exploration differ by about one point on unseen videos, while exploration costs up to 5.7 points during training. On visual pacing, randomized value iteration trails a bandit on per-beat accuracy but earns better whole-video ratings once per-beat feedback is scarce. Against fairly tuned networks, the linear method makes fewer mistakes over 50 videos, although the networks draw level on acceptance.

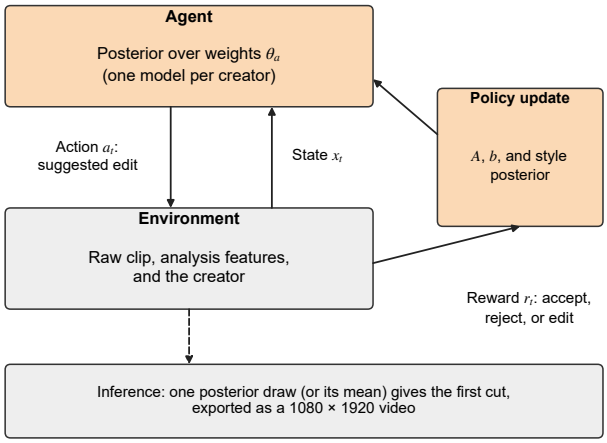


FIGURE 1. Conceptual learning loop. Training and inference share one path: each proposal is applied immediately and every correction updates the model before the next suggestion is scored.

Fig. 1 shows the interaction at the conceptual level. The creator is part of the environment, the state is the feature vector of the current decision, and the policy update happens after each piece of feedback rather than in a separate training phase.

I. RELATED WORK

A. COMPUTATIONAL AND TRANSCRIPT-BASED VIDEO EDITING

A long line of work in graphics and human-computer interaction reduces editing effort by aligning video to its text. Berthouzoz et al. [9] placed cuts and transitions in interview footage from a transcript and hid jump cuts with audio-visual cues. Rubin et al. [10] built content-based tools for audio stories, and Shin et al. [14] linked scripts to audio during authoring. QuickCut [15] matched narration with footage from the editor's own annotations. Video Digests [16] made lecture videos skimmable by segmenting them through their transcripts. B-Script [11] recommended B-roll from the transcript, and Leake et al. [17] encoded film idioms as cost terms for dialogue scenes. Fried et al. [18] went further and

synthesized the talking head itself after a text edit. These systems achieved large time savings, but their preferences are either fixed idioms or settings that the user must supply. None of them learns from one user's corrections across projects.

Recent systems add language models. LAVE [19] lets a user direct edits through an LLM agent. VideoDiff [20] generates alternative rough cuts for comparison. In the VideoDiff evaluation a participant asked for the system to learn personal preferences over time, which is the direct motivation for our work. Write-A-Video [21] assembles montages from text. These tools depend on the user stating an intent in words, which is natural for structural edits and tedious for hundreds of small pacing decisions.

B. REINFORCEMENT LEARNING FOR EDITING

RL has been applied to adjacent tasks. Zhou et al. [22] summarized video with a policy trained on diversity and representativeness rewards. Exposure [6] learned a sequence of interpretable photo operations with deep RL and adversarial rewards. Closest to ours is Hu et al. [23], who trained an RL editor on professional demonstrations with a pretrained vision-language model. These works learn one policy from a dataset for a general notion of quality. They require offline training at scale and do not adapt to an individual. Our agent, in contrast, begins almost empty, trains on one person during normal use, and is evaluated on matching that person.

C. CONTEXTUAL BANDITS AND THOMPSON SAMPLING

Contextual bandits became a standard tool for personalization after LinUCB was applied to news recommendation [3] and analyzed with linear payoffs [12]. Thompson sampling, which dates to 1933 [24], has strong empirical and theoretical support [25]–[28]. Neural contextual bandits [30] replace the linear model with a network, although Riquelme et al. [29] found that simple linear posteriors on good features remain hard to beat. Our results agree with that finding over the first videos, and they also show networks closing the gap as data accumulates. Surveys show the range of applications [31], and a full treatment is given in [32]. For non-stationary preferences, discounted estimators give guarantees [33]. We use these components as they are and focus on the domain-specific structure placed around them.

D. PRIORS ACROSS USERS

A new user's cold start is often handled by sharing information across tasks. Meta-Thompson sampling learns a prior over tasks [34]. Cella et al. [35] bias linear bandits toward a learned mean, and latent bandits [36] assume that each user belongs to one of a few hidden types and infer which. Our style posterior is closest to latent bandits, but its mixture supplies only a prior mean to per-creator ridge models, so the agent can move away from all the prototypes when a creator matches none of them. We test that case explicitly.

E. LEARNING FROM HUMAN FEEDBACK

TAMER [37] showed that agents can learn from human evaluative signals, and later work found that such signals depend on the agent's current policy [38]. Preference-based deep RL [7] and RLHF for language models [8] learn a reward model from human comparisons. Interpreting usage as feedback has a long history in recommendation, where implicit signals are treated as weak evidence [39]. Our reward mapping borrows that caution: a suggestion that was seen but left alone earns half the reward of an explicit accept, and one never seen earns nothing. Choosing what to ask a person is the concern of active learning [40]. Our queue uses the same intuition, with bandit uncertainty as the criterion.

F. PLANNING AND MODEL EXPLOITATION

Posterior sampling for RL (PSRL) [41] extends Thompson sampling to whole episodes, solving an MDP sampled from the posterior. Model-based RL is known to exploit errors in a learned model, which motivated limited rollouts [44] and uncertainty penalties [45]. Constrained MDPs handle budgets formally [46]. Our length-target planner uses PSRL with a known transition model, and we report a concrete failure in which the solver exploited a linear reward model until it cut an entire video. Proximal policy optimization [13] was our first plan, but with known deterministic transitions an exact solution is both cheaper and optimal for the given reward. Where the reward does depend on the state, as with pacing, randomized least-squares value iteration [42], [43] explores coherently across an episode without a separate uncertainty model, and it is the closest sequential relative of the Thompson sampling we already use.

G. OFF-POLICY EVALUATION

Offline evaluation estimates what a different policy would have achieved from logs [47], [48]. Inverse propensity scoring [49], its self-normalized form [50], and doubly robust estimators [51] are the main tools, and slate settings have their own estimators [52]. Adaptive logging complicates inference [53], and bootstraps with few clusters are too narrow [54]. We log propensities for every suggestion so these estimators can be used, and we validate them against simulator truth before applying them to real logs.

H. SPEECH ANALYSIS FOR DISFLUENCIES

The analysis layer relies on modern speech models. Whisper [55] transcribes accurately but tends to tidy disfluent speech, which is why prompting strategies have been studied for verbatim output [56]. Forced alignment with a character model such as wav2vec 2.0 [57] through connectionist temporal classification (CTC) [58] gives accurate word timing, as WhisperX [59] demonstrated. Disfluency detection itself has been studied with sequence models [60], joint parsing [61], and end-to-end systems [62]. These treat removal as a uniform clean-up task, but creators differ on which fillers to keep.

TABLE 1. Qualitative Comparison With Existing Approaches

Approach	Personal	Online	Queue	Length
Silence-removal rules	No	No	No	No
Transcript editors [9], [10]	Manual	No	No	Manual
LLM-directed editors [19], [20]	By prompt	No	No	By prompt
RL from demonstrations [23]	No	No	No	No
Bandit with no style prior (baseline)	Yes	Yes	No	No
Proposed	**Yes**	**Yes**	**Yes**	**Yes**

Personal: adapts to one user. Online: improves during use. Queue: chooses what to ask about. Length: plans to a target duration.

I. WHY A DIFFERENT APPROACH IS NEEDED

Table 1 summarizes these gaps. We found no system that learns one creator's preferences online for this class of edit while also planning to a length target and logging what is needed for counterfactual evaluation, although our search was not exhaustive. The proposed design addresses that combination with models small enough to run per user on a laptop.

II. SOLUTION APPROACH

Our system, RoughCut, turns raw speech footage into a set of editing decisions, proposes an action for each with per-creator agents, and applies all proposals straight away so that the creator starts from a finished first cut. Corrections then flow back through a reward mapping. What sets it apart is what surrounds the bandit: the style posterior used as its prior, the review queue built from its uncertainty, an exact posterior-sampling planner for target lengths, a randomized value-iteration agent for pacing, and propensity logging for counterfactual evaluation. Fig. 2 shows the full pipeline.

A. ENVIRONMENT DESIGN

1) Units and Analysis

Audio is transcribed with faster-whisper running the small English Whisper model [55] at beam size 10, without conditioning on previous text because conditioning caused looping output over silence. The decoder runs in its own process: inside a process that had already prepared the footage, identical audio produced 305 words instead of 324 on one clip and lost a whole sentence. We did not find the internal cause, but three repeated runs in fresh processes were identical. A short disfluency-rich prompt nudges the decoder to keep fillers. Words with no voiced frames beneath them are discarded as hallucinations. Word boundaries are then re-derived by Viterbi alignment of Whisper's characters to the frame-level CTC output of wav2vec 2.0 [57], following WhisperX [59], since Whisper's own timestamps stretch words across pauses.

Two steps go beyond that standard recipe. First, the CTC model's own greedy reading is compared with Whisper's word list using sequence matching, and only two kinds of extra token are trusted: a hesitation sound lasting at least 80 ms, or an exact repeat of a neighboring word. This restored the "I I" stutter in our test clip that Whisper had written as one word. Everything else in the greedy reading is ignored, as its spelling is poor. Second, a filler touching its neighbors is cut only if the audio has a dip at both edges. We search within 80 ms of the boundary for a frame at least 6 dB quieter than the speech on either side and move the boundary onto it. A boundary with neither gap nor dip stays, because cutting there clips a neighboring word. Spoken lines can also be aimed at the wrong audience. A creator may say "this one is for YouTube Shorts" to a colleague, and that line is clean, complete English, so filler and retake detection pass it. Each sentence therefore gets a crew score. Words about production ("voiceover", "one for Instagram") add 0.4 and softer signals add 0.2. Only when such words are present do quiet delivery, isolation between long gaps, and shortness add to the score. Viewer-directed words ("comment", "subscribe") subtract from it, and a score of at least 0.5 cuts the line in the first edit. The creator can restore it like any other cut. These weights were set by hand from very few clips, so we treat the rule as a useful default and not as a validated detector.

The clip is then segmented into four kinds of unit: phrases (split at gaps of at least 0.25 s and at sentence ends, capped near 6 s), pauses (gaps of at least 0.15 s), filler tokens, and caption lines of at most four words. Retakes receive particular care, as they decide which sentences survive. A fuzzy token match catches content retakes. Because most false starts consist of function words ("you'll also be,"), a separate matcher compares all words of a phrase with the speech that follows it, even across phrase boundaries. A complete earlier delivery beats a later repeat only if it is clearly better spoken, with a margin of 0.08 on a score that penalizes fillers and restarts. Phrase embeddings from a sentence transformer [63] provide redundancy and topic-centrality scores, and YuNet [64] tracks the face for a 9:16 crop. All of this runs once, at upload.

2) State Space

Each decision has a context vector $x_t \in \mathbb{R}^d$, normalized and partly binned, with d between 16 and 39 depending on its type (Table 2). Contexts also carry a one-hot video type and a bias term. Decisions are processed left to right, so the context can include history features such as the ratio of output time to source time so far and the time since the last visual change. In the bandit layer these features are the only way an earlier action affects a later state.

3) Action Space and Dynamics

Phrase, pause, filler, pace, and music decisions use disjoint arms, with one weight vector per action. Caption emphasis uses a shared weight vector over per-word features, so what is learned about numbers on one line transfers to all lines. Pause actions leave 0.40 s, 0.15 s, or 0.03 s of silence, or keep the gap untouched. All decisions compile into a single edit decision list, the only representation that both the browser preview and the FFmpeg renderer read. Speech edges are padded by 40 ms, every cut has an 8 ms audio fade, and clip edges snap to the 30 frames per second grid so that the rendered length equals the planned length. The transition

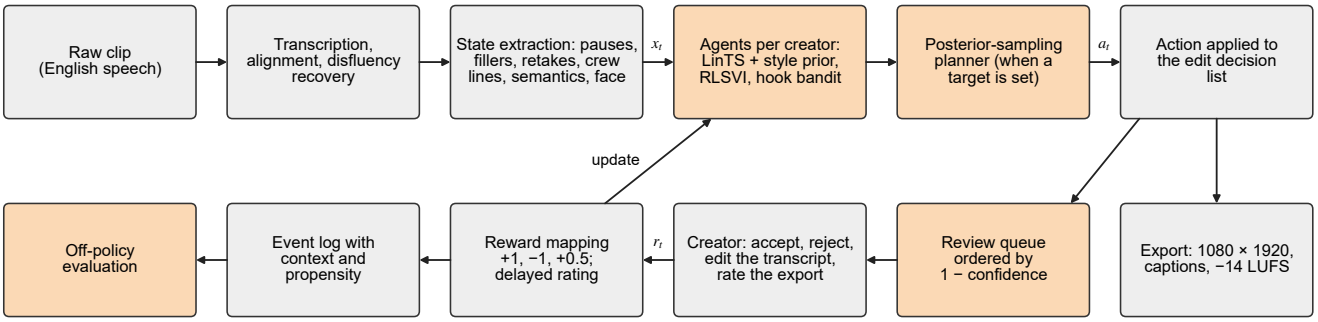


FIGURE 2. Complete pipeline. Orange blocks are the learning components. The top row runs when a clip is uploaded. The bottom row runs each time the creator reacts, and policy state can always be rebuilt by replaying the event log, which is how undo works.

TABLE 2. Decision Types, Actions, and Context Features

Type	Actions	d	Representative features
Phrase	KEEP, CUT	39	duration, retake score, earlier take, incomplete, filler ratio, centrality, redundancy
Pause	KEEP, TIGHT_40, TIGHT_15, REMOVE	21	gap (log and bins), sentence end, voiced ratio, take boundary
Zoom	NONE, Z115, Z130	27	loudness, punctuation, face size, time since last zoom
Filler	KEEP, REMOVE	23	token one-hot, position, gaps, verb-like use of "like"
Emphasis	NONE, or one of up to 3 words	16	per word: number, proper noun, TF-IDF, stress, negation
Pace	S100, S110, S125	18	speech rate, duration, loudness, filler ratio
Music	FROM_START, AFTER_INTRO	7	video length, first-sentence length, type
Beat	NONE, Z115, Z130, TEXT_POP, SFX	29	time since last beat, its kind, streak, centrality
Title	one of up to 3 lines	8	first line, question, number, "you", centrality

dynamics are deterministic and known: applying an action changes only the running output time and history counters. The simulator experiments use a zoom decision with the first three beat actions; the deployed editor replaces it with the beat decision of Section II-C4.

4) Reward Design

Table 3 gives the reward mapping. The reward is defined over the creator's reactions and not over an aesthetic score, so "good" simply means "this creator would keep it". A binary rejection implies that the other arm was right, and a multi-arm rejection gives only negative evidence, which is why exploration could matter there. Manual edits use the same mapping: deleting at least 70 percent of a phrase counts as a phrase override, whereas a smaller word-level cut is applied but not learned from, because it lies outside the action space. Seeing a unit is tracked by the player, so an untouched suggestion only earns the weak implicit reward if the creator could actually have noticed it.

TABLE 3. Reward Mapping for Suggestions and Manual Edits

Creator's reaction	Bandit update
Accept	proposed arm $r = +1$
Reject, binary decision	proposed $r = -1$; other arm $r = +1$
Reject, multi-arm decision	proposed $r = -1$ only
Override to a' (panel, transcript, or timeline)	proposed $r = -1$; a' $r = +1$
Untouched at export, seen	proposed $r = +0.5$
Untouched, never seen	no update
Rating at export, $\rho \in \{-1, 0, 1\}$	rhythm agent only (Section II-C4)
Undo or redo	flip event flag, replay log

5) A Worked Example

A short passage from one of our recordings shows how these parts meet. The speaker says "You'll also be learning about what's the difference between green screen and", stops, and then tries the line several more times ("you'll also be," and so on) before finally delivering "you'll also be learning the difference between green screen and blue screen." The content-token matcher sees little here, as the fragments are almost entirely function words. The all-words matcher finds that each fragment reappears at the start of the later speech, so each one receives `is_earlier_take` $= 1$ and a high retake score. Those two features then enter the phrase context. They do not decide the edit. A learned weight on them does, which is why a creator could in principle teach the agent to keep such a passage. On this recording, the first cut a new creator receives removes every fragment. The pauses around the fragments merge into one cut, so the viewer hears "and so much more." followed directly by the clean take, with an 8 ms fade at the join.

B. RL ALGORITHM

We chose linear Thompson sampling over PPO or deep Q-learning for three reasons. First, feedback is scarce: a creator supplies hundreds of signals, not millions. Second, the linear model's weights can be read back to the creator as an explanation and as a profile. Third, the closed-form posterior makes undo exact, because replaying the log yields the same state. There is no learning rate or batch size to tune, since the

update is a closed-form ridge step applied after every event.

For arm a of a decision type, the expected reward is modeled as $\theta_a^\top x$. Prior and data statistics are held separately:

$$A_a = \lambda I + A_a^d, \qquad \hat{\theta}_a = A_a^{-1}\left(\lambda\,\theta_a^0 + b_a^d\right), \tag{1}$$

$$A_a^d \leftarrow \gamma A_a^d + w\,xx^\top, \qquad b_a^d \leftarrow \gamma b_a^d + w\,r\,x, \tag{2}$$

where λ is the prior strength, θ_a^0 the prior mean, γ the discount, and w an event weight. To act, the agent draws $\tilde{\theta}_a \sim \mathcal{N}(\hat{\theta}_a, v^2 A_a^{-1})$ and takes

$$a_t = \arg\max_a \tilde{\theta}_a^\top x_t, \tag{3}$$

replaced by a uniform arm with probability ϵ. The action probabilities are estimated from $M = 200$ draws,

$$p(a \mid x) = (1-\epsilon)\,\frac{1}{M}\sum_{m=1}^{M} \mathbf{1}\left[a = \arg\max_{a'} \tilde{\theta}_{a'}^{(m)\top} x\right] + \frac{\epsilon}{K}, \tag{4}$$

and $p(a_t \mid x_t)$ is stored as both the logged propensity and the confidence of the suggestion. The goal is to minimize regret against the creator's own ideal action a_t^*,

$$R_T = \sum_{t=1}^{T} \mathbf{1}\left[a_t \neq a_t^*\right], \tag{5}$$

which counts wrong proposals and is what a creator pays in corrections.

Replaying the log gives the same posterior as the online run because every path through the system, whether from the simulator, the application programming interface, or an undo, applies events through a single learner object. We test that property directly. It also means that a change of prior or a schema migration does not throw away what a creator has taught the agent, since the raw context is stored with each suggestion and the sufficient statistics can be rebuilt.

Linear models also allow an explanation of each suggestion. For the proposed action and its runner-up, the difference of scores decomposes into per-feature contributions $(\hat{\theta}_a - \hat{\theta}_{a'})_j\,x_j$. The three largest are turned into a sentence in the "Why?" panel, and the largest weights across all decisions form the creator's profile page. These are statements about the model and not about the creator's actual reasoning, and we have not studied whether people find them useful.

C. NOVEL COMPONENTS

1) Tempered Style Posterior as Prior

For K style prototypes $\theta^{(k)}$, ridge-fitted offline, and a pooled generic prior, every event updates a log weight,

$$\log w_k \leftarrow \log w_k - \frac{\left(r - \theta_a^{(k)\top} x\right)^2}{2\sigma^2}. \tag{6}$$

With $\sigma = 0.7$, the posterior concentrates after a few events and drives the sentence shown to the creator. The bandit, however, uses weights tempered by $(\sigma/\sigma_p)^2$ with $\sigma_p = 3$ to form its prior mean $\theta^0 = \sum_k \tilde{w}_k \theta^{(k)}$. Tempering separates two quantities that a plain latent-bandit posterior ties together: what we say about the creator and how far one click can move the policy. Without it, one rejected KEEP flipped 65 percent of kept phrases on the next video at $\lambda = 10$. With a stronger keep/cut prior of $\lambda = 30$ that figure fell to 8 percent. The same safety argument removes the ϵ floor on keep/cut in the product, since a uniform draw there can delete a sentence, and Thompson draws still explore where the posterior is unsure.

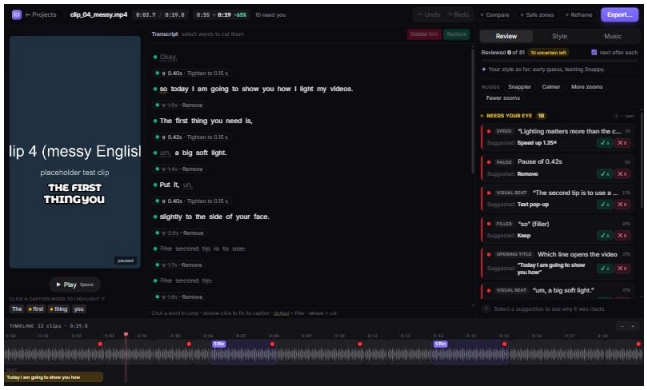


FIGURE 3. Editor on a synthetic test clip: captioned preview, transcript with struck-out cuts, review queue sorted by uncertainty, and a timeline with confidence dots.

2) Uncertainty-Ordered Review

Suggestions are sorted by $1 - p(a_t \mid x_t)$, and those with confidence below 0.7 form the "needs your eye" list. Everything else stays applied and available for a spot check. After each correction, unreviewed suggestions are re-scored, with two guards: an action is changed only if Thompson sampling would pick it at least 80 percent of the time, and one nudge moves at most a quarter of a decision type. Fig. 3 shows the resulting interface.

3) Target-Length Planning by Posterior Sampling

A duration target cannot be enforced one decision at a time. We define an MDP over the phrase and pause decisions in time order. Its state is $s = (o, c)$, with output seconds o on a 0.1 s grid and a flag c for whether the previous phrase was kept. The reward at stage k is the learned model, and the terminal value penalizes missing the target T:

$$r_k(a) = \theta_a^\top x_k, \tag{7}$$

$$V_K(o) = -\beta \max\left(0, \tfrac{o}{T} - 1 - \tau\right) - \beta_u \max\left(0, 1 - \tau - \tfrac{o}{T}\right), \tag{8}$$

with $\beta = 150$, $\beta_u = 6$, and $\tau = 0.05$, so a target works mainly as a limit and a clip whose good edit is shorter is not padded. Backward induction,

$$V_k(s) = \max_a \left[r_k(a) + V_{k+1}\left(s'(s,a)\right)\right], \tag{9}$$

gives the exact optimum, because transitions are known. A phrase that continues the sentence of the previous phrase must share its fate, so a length target can never strand half a sentence.

Exploration follows PSRL [41]: once per video, θ is drawn from the bandit posterior and the MDP is solved for that draw. The plan is closed loop, so a creator's override simply moves the state to another row of the solved table. Propensities are the share of 16 further posterior draws whose optimal plan takes the same action at that state. Since the planner's step model only approximates the rendered edit, the planned target is corrected against the real edit length, up to twice per video.

The design decision worth reporting is the treatment of history features. In the first version, they followed the MDP state. The result looked exact but was badly wrong: on simulated test videos the solver found that cutting every phrase was optimal, as the linear model had never seen a state like "everything cut so far" and extrapolated high reward there. This is the model exploitation problem described in model-based RL [44], [45], here in a miniature and exactly solvable form. We now hold those features at their values along the bandit's own posterior-mean trajectory, which keeps the reward inside the data's support. The cost is that the reward no longer reacts to the plan's own history, so the only sequential coupling that remains is the length budget and the sentence ties.

The cost of an exact solution is modest. For K stages, n buckets of output time, and at most four actions, backward induction takes $O(K\,n)$ table operations per solve, vectorized over the grid. A typical video has about 65 stages and a few hundred buckets. One complete planning call involves up to three solves for calibration and 17 more for the acting plan and the propensity draws, and it finishes in about 0.2 s on a laptop CPU (Table 9). A sampled policy-gradient method would need far more computation to approximate the same table, and its result would be harder to explain.

4) Visual Rhythm by Randomized Value Iteration

Whether a punch-in works depends on what came before it, so this decision is sequential by nature. At each kept phrase longer than 1 s, the agent chooses among no effect, a 1.15x or 1.3x zoom, a text pop-up, and a sound effect. The state $\phi(s) \in \mathbb{R}^{29}$ holds binned time since the last beat, the kind of that beat, the streak length, beats per minute so far, phrase features, and products of time bins with topic centrality, because the interesting rules are conjunctions. Unlike the planner's frozen features, these are true state: the chosen action changes the next ϕ.

With a linear $Q(s,a) = \theta_a^\top \phi(s)$, RLSVI [42] perturbs each reward with Gaussian noise and each prior with a draw, and then runs fitted value iteration on the perturbed data,

$$y_i = r_i + \xi_i + \gamma \max_{a'} \theta_{a'}^\top \phi(s_i'), \qquad \xi_i \sim \mathcal{N}(0, \sigma^2 / w_i), \tag{10}$$

$$\theta_a \leftarrow \Big(\sum_{i:a_i=a} w_i \phi_i \phi_i^\top + \lambda I \Big)^{-1} \Big(\sum_{i:a_i=a} w_i \phi_i y_i + \lambda \tilde{\theta}_a^0 \Big), \tag{11}$$

with $\gamma = 0.9$, $\sigma = 0.15$, $\lambda = 1$, and 10 sweeps. The agent then acts greedily on that one draw for the whole video, which is Thompson sampling extended to an episode. The prior mean θ^0 is fitted to a calm hand-written rule, one modest beat after several quiet seconds. The rating at export is delayed feedback: it adds 2ρ to the last step, for value iteration to carry backward, and 0.5ρ to every step. Those two weights were chosen on tuning seeds only. Algorithm 1 gives the procedure, and setting $\gamma = 0$ turns the same code into a bandit on the same features, which is the baseline we compare with.

Algorithm 1 Visual Rhythm Agent (RLSVI)

Require: Prior θ^0, logged episodes $\mathcal{D}$
for each new video **do**
  Perturb prior: $\tilde{\theta}^0 \leftarrow \theta^0 + \sigma \lambda^{-1/2} \varepsilon$
  Perturb rewards in $\mathcal{D}$; add the rating to steps
  for 10 sweeps **do**
    Compute targets by (10); update θ_a by (11)
  end for
  for each kept phrase in time order **do**
    Observe $\phi(s)$; take $a = \arg\max_a \theta_a^\top \phi(s)$
    Observe the creator's reaction; advance the state
  end for
  Observe the rating ρ; append the episode to $\mathcal{D}$
end for

5) Opening Title, Pace, and Music Timing

Three smaller decisions use the same machinery. A hook bandit picks the on-screen title from up to three candidate lines, using a shared Thompson model over line features, so a preference for questions carries across videos. It does not move footage to the front. Phrase speed and music timing are ordinary LinTS decisions with a fixed bonus on the plain option, so an untrained agent stays at normal speed. Cosmetic decisions are capped in the review queue, since they are numerous and cheap to get wrong. We do not evaluate these three decisions in the simulator, so they are product features here and not research results.

6) Counterfactual Evaluation

Each logged suggestion stores (x_i, a_i, p_i, r_i). For a target policy π, with weight $w_i = \pi(a_i \mid x_i)/p_i$ and a cross-fitted ridge reward model $\hat{q}$, we compute the direct method (DM), inverse propensity scoring (IPS), self-normalized IPS (SNIPS), and doubly robust (DR) estimates, plus self-normalized DR (SNDR):

$$\hat{V}_{\mathrm{DR}} = \frac{1}{n} \sum_{i=1}^{n} \Big[\sum_a \pi(a \mid x_i) \hat{q}(x_i, a) + w_i \big(r_i - \hat{q}(x_i, a_i) \big) \Big]. \tag{12}$$

Decisions within a video are dependent, so the bootstrap resamples whole videos. With few videos, percentile intervals were too narrow [54]. We therefore report the estimate $\pm\, t_{G-1,0.975} \sqrt{G/(G-1)}$ times the bootstrap standard deviation, for G videos.

Algorithm 2 Online Editing Agent for One Creator

Load prototypes; set $\log w_k \leftarrow 0$, $A^d, b^d \leftarrow 0$
Set prior mean θ^0 from tempered w
for each uploaded video **do**
  Analyze the clip into units and features
  if a target length is set **then**
    Solve the planner (Algorithm 3)
  end if
  for each decision t in time order **do**
    Observe context x_t (static features and history)
    $a_t \leftarrow$ planner action if planned, otherwise (3)
    Compute and log $p(a_t \mid x_t)$ by (4); apply a_t
  end for
  Sort the review queue by $1 - p(a_t \mid x_t)$
  while the creator is editing **do**
    Observe an event; map it to updates (Table 3)
    Update A^d, b^d by (2) and w by (6)
    Recompute θ^0; re-score unreviewed suggestions
    On undo: flag the event and replay the log
  end while
  At export: give seen, untouched suggestions $r = +0.5$
end for

Algorithm 3 Target-Length Planner with Posterior Sampling

Require: Stages with frozen history features, target T, posterior
Solve (9) with posterior means $\hat{\theta}$
for up to two calibration rounds **do**
  Roll out the plan and measure the real edit length L
  if $|L - T|/T \leq \tau/2$ **then** stop
  Shift the planned target by $T - L$ and solve again
end for
Draw $\tilde{\theta} \sim$ posterior; solve (9): acting plan
Solve for 16 more draws: propensity plans
for each stage reached, at the state (o, c) actually observed **do**
  Act by table look-up in the acting plan
  Propensity $\leftarrow$ share of draws with the same action
end for

D. ALGORITHM DESCRIPTION

Algorithm 2 gives the online loop, Algorithm 3 the planner, and Algorithm 1 the rhythm agent. There is no stopping condition during use, because the agent learns for as long as the creator keeps editing, with style weights frozen after 200 events. Each simulated experiment stops after a fixed number of videos.

E. SYSTEM IMPLEMENTATION

The method is deployed in a complete editor so that feedback comes from realistic use. The backend is written with FastAPI and stores projects, suggestions, and events in SQLite. The frontend, in React and TypeScript, renders the edit decision list in the browser with skip ranges, CSS transforms for crop and zoom, and HTML captions, so no server render is needed while editing. Export uses an FFmpeg filter graph that crops from the original source, burns in captions, applies high-pass filtering, and normalizes loudness to -14 LUFS over two passes. Optional background music is ducked under speech with sidechain compression.

Batch uploads use the posterior mean without exploration. A video is marked ready to post when the sum of $1-$confidence over its suggestions is below 3 and the creator's recent acceptance rate is at least 0.85. One-click export from that state deliberately produces no training signal, because nobody reviewed anything. The 317 automated backend tests cover the reward mapping, replay, and timing rules, and a further 166 cover the interface. They verify the implementation and do not validate the method.

TABLE 4. Simulated Creators (Hidden Ground Truth)

Persona	Pauses	Phrases, fillers, zoom
Snappy	remove over 0.5 s, else 0.15 s	cut fillery or garbled; remove all fillers; zoom on loud lines
Relaxed	keep under 0.8 s, else 0.40 s	cut earlier takes; remove um, uh; no zoom
Storyteller	0.40 s after sentence end, else remove	keep story; remove um, uh, you know; zoom on ? and !
Minimalist	remove all	keep most central up to 30 s; no zoom
Focused	all to 0.15 s	cut off-topic and repeats; zoom on central lines
Blend	as Snappy	phrases and zooms as Storyteller

III. EXPERIMENTAL RESULTS

A. EXPERIMENTAL SETUP

1) Environment and Data

Offline RL datasets of real creators' corrections do not exist for this task and we have not yet collected our own, so the main evaluation uses a simulator. Synthetic videos are generated as timed transcripts with planted mistakes: abandoned takes followed by retakes, tangents, fillers, long pauses, and joins between takes. They then pass through the same segmentation, filler, pause, and feature code as real footage, so the agent sees the same state representation. Each video has roughly 160 decisions.

Simulated creators are six personas with hidden rules for every decision type (Table 4). A persona reviews a proposal, gives a label that is wrong with probability 0.05, and corrects a disagreement with probability 0.6. Half of these corrections are made as manual edits passing through the real classification code. All priors are fitted without the persona under test (leave-one-style-out), on training videos disjoint from the evaluation videos, so a persona's own prototype is never available to it. We also use three recorded clips for the analysis layer.

2) Hardware and Software

All runs used one Windows 11 laptop with an Intel processor reporting 16 logical cores and no GPU. The software was Python 3.10, NumPy 2.2, SciPy 1.15, faster-whisper 1.2, PyTorch 2.14 (CPU), Transformers 5.18, Sentence-Transformers 6.1, OpenCV 5.0, and FFmpeg 9.0. Agents are our own implementation, and the deep baselines are written directly in PyTorch. Gymnasium and Stable Baselines3 are not used, because the environment runs the product's own decision code.

TABLE 5. Hyperparameters

Parameter	Simulator	Product
Prior strength λ	2	10 (keep/cut 30)
Posterior scale v	0.25	0.15
Discount γ	0.995	0.995
Exploration floor ϵ	0.02	0.005 (keep/cut 0)
Thompson draws M	200	200
Style width σ / prior width σ_p	0.7 / none	0.7 / 3
Implicit reward	0.5	0.5
Style frozen after	200 events	200 events
Planner β, β_u, τ, grid, draws	150, 6, 0.05, 0.1 s, 16	
Episodes per run / seeds per persona	10 / 10	
RLSVI γ, σ, λ, sweeps	0.9, 0.15, 1, 10	
Retuned λ, v, ϵ (Section III-D1)	2, 0.15, 0	
Neural: hidden, batch, lr, ϵ, steps per event	64–64, 32, 3×10^{-4}, 0, 4	
NeuralLinear: lr, v, refit interval	10^{-3}, 0.1, 32	

3) Hyperparameters

Table 5 lists the settings. For the linear agents, the terms learning rate and batch size have no direct counterpart. The update of (2) is closed-form and applies per event, so the closest analogs are λ and γ. An episode is one video and an environment step is one decision. Unless noted, results use ten seeds per persona, or 60 runs per method, and ten videos per run. Intervals are 95 percent t intervals over runs.

B. PERFORMANCE EVALUATION

1) Metrics

Acceptance is the share of proposals equal to the creator's ideal action. Return per decision is the mean realized reward from Table 3. We use cumulative regret from (5). Success, the task-level condition, means that a video reaches at least 85 percent acceptance, measured over videos 6 to 10. Steps to target is the number of decisions until a run first completes a video at or above 85 percent. Corrections per minute of footage measure the effort that remains. For target lengths we report relative length error measured on the real edit, and for review the share of correct decisions in the exported edit. Together these are the problem-specific metrics for this domain.

2) Learning and Regret

Fig. 4 shows return and regret over environment steps. Return rises steeply over the first 500 decisions, about three videos, and then flattens. Both fixed baselines stay flat, as expected, and the learners separate from them on the first video. Table 6 gives the numbers. The proposed method ends at 91.0 percent acceptance with 181.5 wrong proposals in 10 videos, whereas Rule makes 668.9 and Global 780.1.

3) Acceptance by Decision Type

Fig. 5 separates the five decision types. Pauses are where rules do worst, at 0.42 on video 10, because the personas disagree most on them, and they are where learning gains most (0.96). Keep/cut starts high for every method because most phrases should be kept, and here the Rule baseline's 0.87 is close to the learners. Emphasis is the hardest learned decision at 0.88, as our features only partly explain the choice of word.

Table 7 repeats the comparison per creator. The proposed method is ahead of Rule for all six, by 19 to 50 points. The gap is smallest for Blend and Snappy, whose pause rule resembles the one in the Rule baseline, and largest for Relaxed and Minimalist, whom a fixed cut-everything-over-half-a-second policy serves poorly. Against plain LinTS the style prior helps most on Snappy and Blend. On Minimalist and Focused it is slightly behind (0.857 against 0.870, and 0.915 against 0.929), so the mean improvement is not a uniform one.

Table 8 gives the same breakdown by decision type. The product settings, which use a stronger keep/cut prior, trade about three points of keep/cut acceptance (0.876 against 0.905) for protection against sudden large changes, and they recover the difference on pauses and zooms. This is a deliberate choice. A wrong CUT removes what someone said, whereas a wrong pause length is hardly noticed, so we accept lower simulated keep/cut acceptance in exchange for slower movement on that one decision type.

4) Stability Across Seeds

The standard deviation of per-seed mean acceptance at video 10 is 0.013 for the proposed method and 0.018 for the plain bandit, both small compared with the 0.32 gap from Rule. Most of the spread across all runs (0.045) comes from differences among personas and not from randomness, since Relaxed reaches 0.974 and Minimalist only 0.857. Minimalist's rule depends on a per-video ranking of centrality that a linear model cannot fully represent.

5) Computational Efficiency

Table 9 shows the learner's cost, measured on one thread. A suggestion takes well under 1 ms on average, so it is the media analysis and not the learning that determines waiting time. The bandit has 305 parameters per creator and its data statistics occupy 67 KiB. The rhythm agent adds 145 parameters and takes 16 ms to prepare a video after ten videos of history. A neural bandit run with one gradient step per event took 20 s against 2.4 s for LinTS. Analysis took 98 s for the 54.5 s clip, 119 s for the 44.2 s clip, and 270 s for the 238.5 s clip, measured while the test suite was running on the same machine, so these are upper figures.

C. RESULTS AND DISCUSSION

1) Does the Style Prior Fix the Cold Start?

On videos 1 to 3, acceptance is 0.824 with no prior, 0.831 with a generic prior, and 0.845 with the style mixture. The improvement is real but modest, and it fades as expected: by video 6 the three curves lie within one point of each other (0.889, 0.889, 0.899). Personal data eventually dominates any prior. What the prior buys is time. The median number of decisions to reach an 85 percent video drops from 430 to 344, about half a video sooner. The style itself is identified (weight above 0.8) after a median of 3.5 events, although

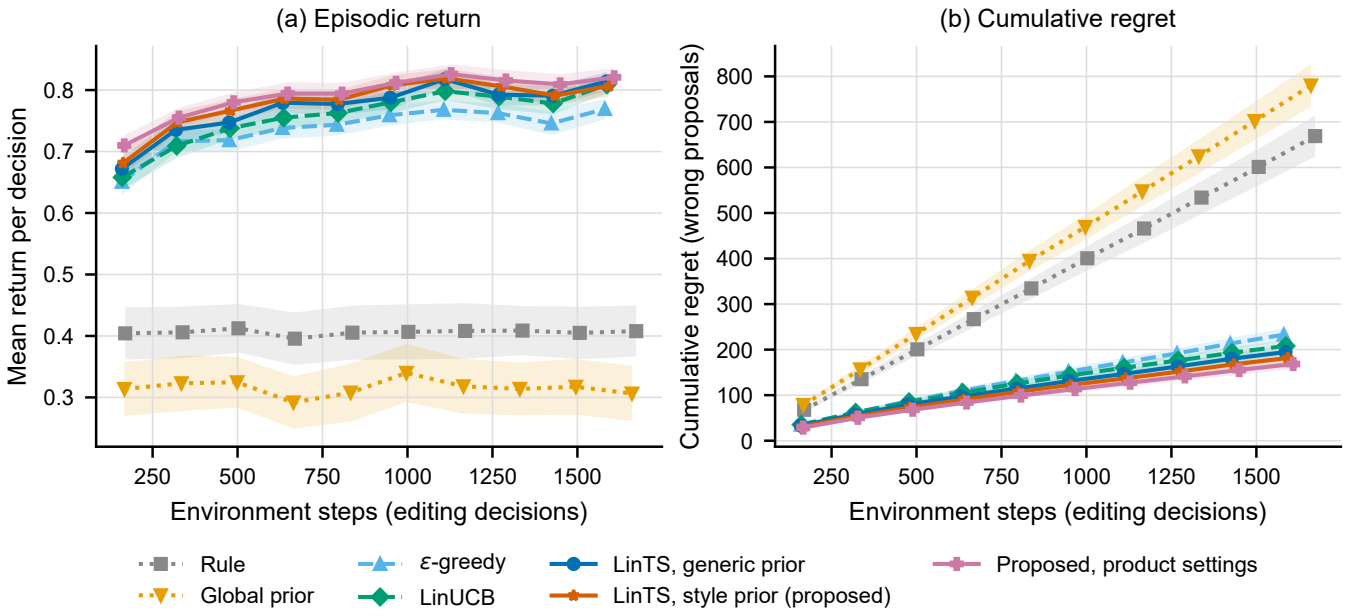


FIGURE 4. (a) Mean return per decision and (b) cumulative regret over environment steps, across six creators and ten seeds. Bands show 95 percent confidence intervals. The learners rise quickly and stabilize, and their bands narrow, whereas the fixed policies stay flat.

TABLE 6. Comparison With Benchmark Methods (Six Creators, Ten Seeds, Ten Videos; ↑ Higher Is Better, ↓ Lower Is Better)

Method	Video 1 ↑	Video 5 ↑	Video 10 ↑	Return ↑	Success ↑	Steps to 85% ↓	Regret ↓	Corr./min ↓
Rule (one-click tool)	0.594	0.596	0.586 ± 0.031	0.408	0.000	never	668.9 ± 45.0	19.91
Global prior, no learning	0.525	0.513	0.521 ± 0.035	0.306	0.000	never	780.1 ± 46.5	22.77
ϵ-greedy	0.777	0.862	0.874 ± 0.011	0.769	0.700	519.0	232.8 ± 14.1	7.64
LinUCB	0.779	0.873	0.903 ± 0.011	0.809	0.823	517.5	208.1 ± 13.0	6.66
Neural, blank	0.735	0.894	0.915 ± 0.011	–	0.880	402.5	186.7 ± 13.6	–
Neural, pretrained	0.789	0.900	0.921 ± 0.011	–	0.947	341.5	167.8 ± 14.2	–
NeuralLinear, blank	0.677	0.880	0.907 ± 0.013	–	0.883	528.0	219.6 ± 15.0	–
NeuralLinear, pretrained	0.724	0.880	0.916 ± 0.012	–	0.857	481.0	202.3 ± 13.6	–
LinTS, generic prior	0.790	0.891	0.904 ± 0.011	0.814	0.853	430.0	194.7 ± 14.7	6.27
LinTS, style prior (proposed)	0.806	0.890	0.910 ± 0.012	0.807	0.903	344.0	181.5 ± 11.0	6.19
Proposed, retuned	**0.826**	**0.909**	**0.931 ± 0.008**	–	**0.967**	**334.0**	**152.7 ± 10.6**	–
Proposed, product settings	0.825	0.900	0.912 ± 0.012	0.821	0.907	343.5	168.4 ± 11.7	5.82

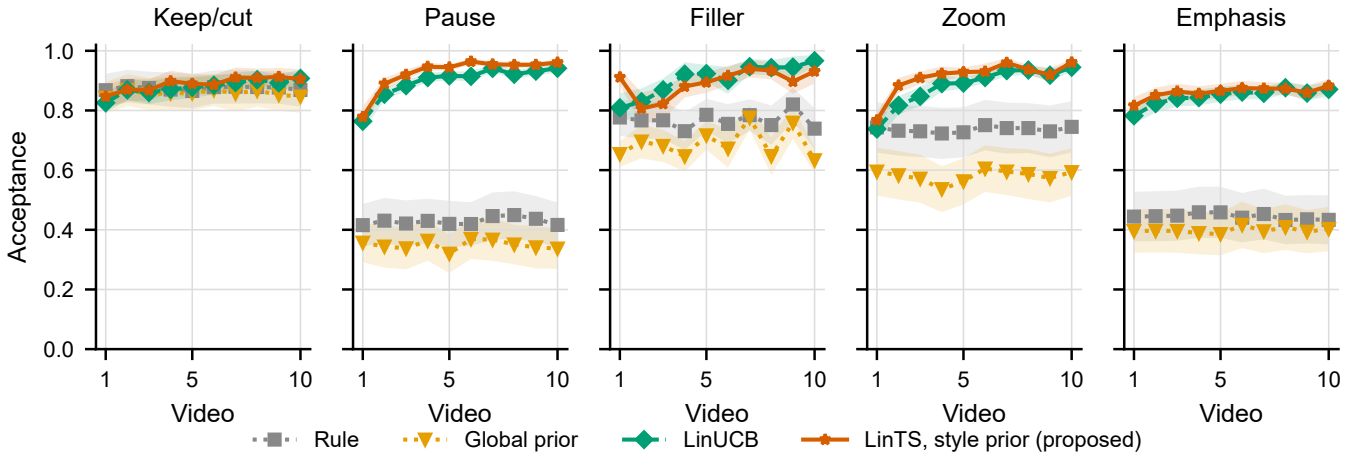


FIGURE 5. Acceptance per decision type over ten videos. The gain from personalization is largest on pauses, which is where creators' preferences differ the most. On keep/cut, a good rule is already close to the learners.

TABLE 7. Acceptance on the Tenth Video per Simulated Creator

Creator	Rule	Global	LinUCB	LinTS	Proposed
Snappy	0.701	0.644	0.881	0.868	**0.893**
Relaxed	0.472	0.348	0.956	0.961	**0.974**
Storyteller	0.671	0.600	**0.912**	0.906	0.910
Minimalist	0.422	0.346	0.860	**0.870**	0.857
Focused	0.531	0.535	**0.933**	0.929	0.915
Blend	0.721	0.652	0.875	0.887	**0.910**

TABLE 8. Acceptance on the Tenth Video per Decision Type

Method	Keep/cut	Pause	Filler	Zoom	Emphasis
Rule	0.869	0.416	0.739	0.745	0.433
Global prior	0.847	0.337	0.633	0.593	0.403
LinUCB	0.907	0.941	0.967	0.945	0.871
LinTS, generic prior	0.904	0.953	0.931	0.963	0.868
Proposed	0.905	0.961	0.931	0.961	0.883
Product settings	0.876	0.983	0.943	0.973	0.889

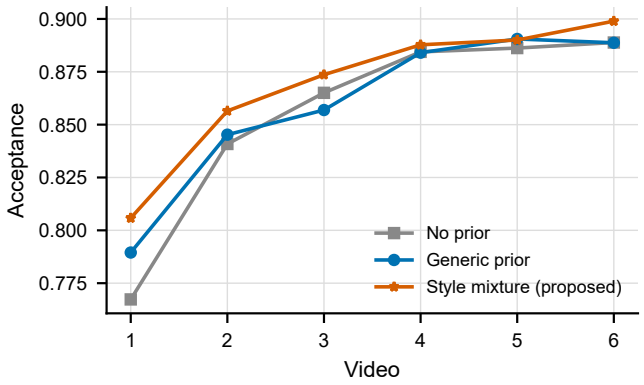


FIGURE 6. Acceptance on the first six videos under three priors, with the evaluated style excluded from the prior. The style mixture leads early and the curves meet as personal data accumulates.

Snappy takes 7.5 and at most 21, as it overlaps with Blend. These numbers come from leave-one-style-out priors, so they probably describe a new creator somewhat similar to the known styles, and a person who resembles none of them should expect less.

Fig. 6 shows how the gap closes. The lead of the style mixture is largest on the first video (0.806 against 0.767 with no prior) and is nearly gone by the fourth, where all three are near 0.885. A paired test on the first three videos gives a gain of 1.5 points over the generic prior ($p = 0.004$). In practical terms, this is about 13 fewer wrong proposals across ten videos against the generic prior, and it costs almost nothing to compute: one mixture update per event.

2) Is the Policy Personalized?

Fig. 7 trains on one persona and evaluates the frozen policy on each. The diagonal averages 0.936 and the off-diagonal 0.468, so a policy that is excellent for one creator is close to useless for another. Blend and Snappy are the expected exception at about 0.85 to 0.87, since they share pause, filler, and emphasis rules by construction. This result supports the premise that a single global model cannot serve everyone and

TABLE 9. Computational Cost of the Learner (One Thread)

Operation	Mean	95th percentile
Thompson choice with 200 draws (ms)	0.07	0.10
Update of one arm (ms)	0.02	0.03
Plan one video, 65 stages (ms)	186	197
Parameters per creator	305	
Data statistics per creator (KiB)	67	
Simulated run, 10 videos, 1600 decisions (s)	2.4	

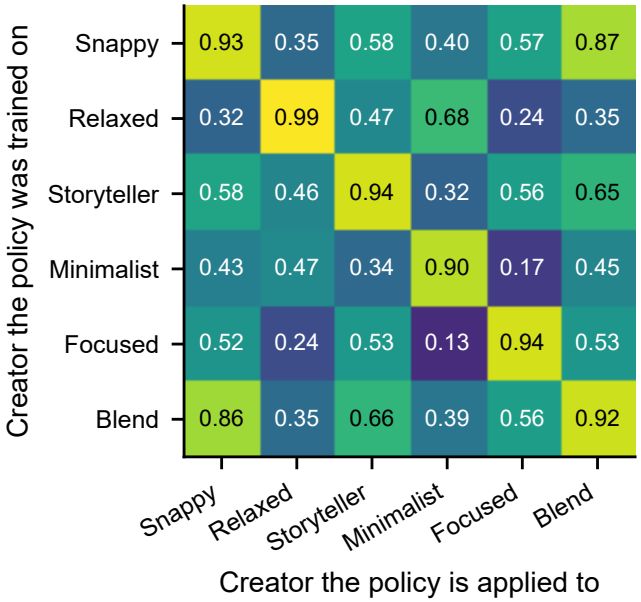


FIGURE 7. Frozen policies applied to each creator on unseen videos. The dominant diagonal shows that learned policies are specific to the creator they were trained on.

is the largest effect in the paper. It should still be read with care, because we wrote personas to differ.

3) Can It Follow Changing Taste?

Fig. 8 shows a creator who switches from Relaxed to Snappy at video 6. Acceptance collapses to 0.54 with no discounting, 0.61 with $\gamma = 0.995$, and 0.68 with $\gamma = 0.98$, and all three recover above 0.8 within three to four videos. Stronger discounting softens the immediate drop but pays for it, with lower acceptance before the switch (0.935 against 0.950) and at the end of the run (0.834 against 0.871). Discounting costs accuracy in stationary conditions, as the ablation also shows, so the choice of $\gamma = 0.995$ is a compromise rather than an optimum.

4) Does Length Need Planning?

Fig. 9 compares the bandit and the bandit with planner, both trained online and then frozen on four unseen videos, for targets of 20, 30, and 45 s. The bandit's length error remains 40.0 percent after six videos, because nothing in a per-decision policy knows the target. With the planner, the error ends at 4.6 percent, below the 5 percent tolerance, and it is small from the first video onward because the target is

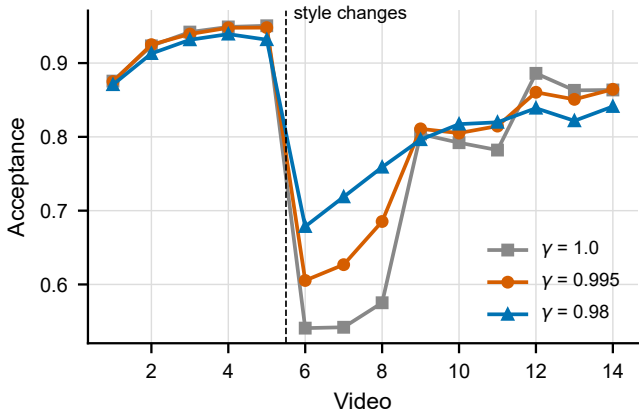


FIGURE 8. Style change at video 6 for three discount factors. Stronger discounting softens the drop at the switch but lowers the level reached before and after it.

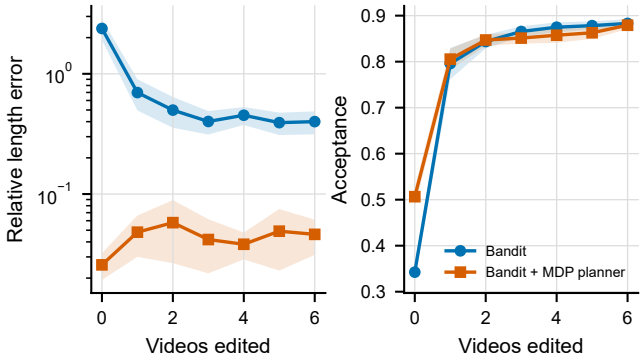


FIGURE 9. Unseen videos for a creator with a target length. Left: relative length error on a logarithmic axis. Right: acceptance.

enforced by planning rather than learned. It costs 0.4 points of acceptance (0.879 against 0.883), which lies within the intervals. Keep/cut acceptance stays near 0.68 for both, so neither reproduces Minimalist's choice of which phrases to keep. Planning fixes the length without fixing the selection, and this is the method's clearest unsolved problem.

Table 10 splits the result by target. Tighter targets are harder for the bandit, with 54.7 percent error at 20 s, as a 20 s target needs more cutting than any learned habit provides, whereas the planner stays under 6 percent at all three.

5) Does the Queue Save Effort?

Fig. 10 varies the share of decisions a creator is willing to review. At 30 percent, an uncertainty-ordered queue gives 0.869 correct decisions in the export, against 0.784 for random order and 0.633 for timeline order. At full review, all three are the same as they must be, at about 0.90. The gain comes from catching errors: confidence predicts mistakes well enough that the same effort fixes more of them. Policies trained at a 30 percent budget also score higher when frozen after uncertainty-ordered review (0.841 against 0.800), which suggests that the queue selects useful training samples too, though the difference is near the width of the intervals.

TABLE 10. Relative Length Error on Unseen Videos by Target

Method	20 s	30 s	45 s
Bandit only	0.547	0.340	0.314
Bandit with planner	0.054	0.052	0.033

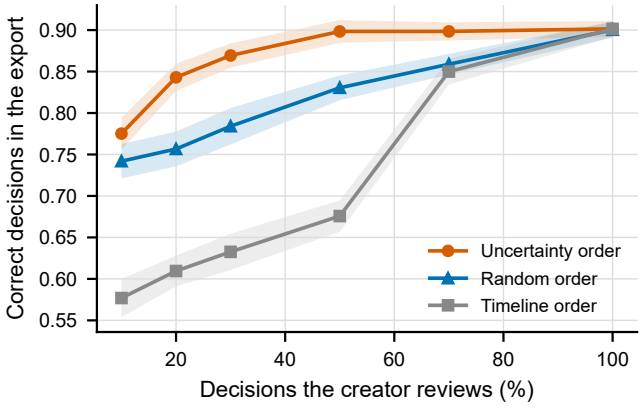


FIGURE 10. Quality of the exported edit against review budget. Uncertainty order dominates below full review and converges with the others at 100 percent.

6) Can Confidence Be Trusted?

The queue only works if confidence tracks mistakes, so we measured that directly (Table 11). At $v = 0.25$ the expected calibration error over about 58 000 suggestions is 0.023. Suggestions below the 0.7 threshold make up 10.1 percent of the total and are right only 37.1 percent of the time, whereas those above it are right 94.4 percent of the time. The area under the curve for telling right from wrong is 0.89. So a creator who looks only at the flagged tenth is looking where most errors are.

This is what a smaller v loses. At $v = 0.05$, acceptance is higher (0.916 against 0.902 late in training), but the area under the curve falls to 0.72 and only 3.2 percent of suggestions fall below the threshold, so many mistakes go unflagged. At $v = 0.5$ the ranking is sharper (0.94), but 18.6 percent are flagged and acceptance falls to 0.876. The experiments use $v = 0.25$ and the product $v = 0.15$, chosen between those extremes: on 120 fresh runs the product's values reached 0.903 against 0.892 after six videos (paired difference 0.011 $\pm$ 0.005). This result also gives a more accurate account of what Thompson sampling contributes in this system. It does not seem to help learning under our feedback model, but a posterior of the right width produces an uncertainty signal that a creator can use.

7) Can a Policy Be Judged Offline?

Fig. 11 and Table 12 evaluate policies from LinTS logs, with simulator truth as reference. On average the doubly robust estimate ranks the learned policy first by a wide margin and comes within 1.3 points of it (true 0.932, DR 0.945), and it recovers the order of Rule (0.749, DR 0.739) and the generic prior (0.703, DR 0.694). DM has the lowest error (RMSE 0.065) but its bias makes its intervals almost useless,

TABLE 11. Confidence as a Predictor of Mistakes

Measure	$v = 0.05$	$v = 0.25$	$v = 0.5$
Expected calibration error	0.065	0.023	0.023
Area under the curve (right vs. wrong)	0.715	0.893	0.940
Share flagged (confidence < 0.7)	0.032	0.101	0.186
Correct among flagged	0.296	0.371	0.372
Correct among not flagged	0.919	0.944	0.963
Late acceptance (videos 8–10)	0.916	0.902	0.876

TABLE 12. Off-Policy Estimators Against Simulator Truth

Estimator	Bias	RMSE	95% CI coverage	Picks best
DM	0.040	0.065	0.48	1.00
IPS	−0.007	0.153	0.72	0.72
SNIPS	0.013	0.124	0.81	0.97
DR	0.002	0.126	0.91	0.94
SNDR	0.003	0.121	0.86	0.94

covering the truth only 48 percent of the time. DR has much smaller bias (0.002) and the best coverage, 91 percent. The small-sample t correction raised coverage from the 71 to 78 percent we saw with percentile intervals, but it still falls short of the nominal 95 percent. Part of the shortfall is likely adaptive logging [53]. The effective sample size is just 15 percent of the 950 logged decisions per creator, because LinTS quickly becomes confident and stops supporting other actions. Offline evaluation in this setting therefore works for ranking clearly different policies, but not for detecting small differences.

8) Real Audio

Simulation says nothing about whether the speech analysis works, so we ran three clips through the actual pipeline and compared the first cut a brand-new creator would receive with a reference edit (Table 13). The first is a 44.2 s recording of a person with a real chain of restarts: four abandoned attempts at one sentence followed by the clean take, and unrelated background speech in a 7 s gap. The first cut drops every abandoned attempt. The second is a 238.5 s recording with long gaps, six retaken lines, and five phrases spoken to someone off camera. All five were cut and the call to action at the end was kept. The third is a 54.5 s scripted stress clip with a synthetic voice, containing fillers, a stutter, hedges, a three-attempt restart, and a sentence where the first delivery is the good one. The agent resolved all of them and retained one filler ("basically"), which has no audio dip around it.

These are encouraging results but very thin evidence. The sample is three clips and one of them has a synthetic voice. The references were prepared with AI assistance from the transcript and from decoding doubtful stretches on their own, not by an independent editor listening, and they use the transcript's wording, so zero differences does not mean perfect transcription: the system writes "Klein" for "client", and no reference corrects it. The crew rules and the process isolation were developed on these same clips, so they are not held out. Through this comparison we found and fixed a stuttered scrap of a restart and the lost sentence described

TABLE 13. First Cut Against a Reference Edit on Recorded Clips

Measure	Clip A	Clip B	Stress clip
Source length (s)	44.2	238.5	54.5
First cut length (s)	22.1	72.6	19.2
Clips in the edit	6	24	12
Suggestions / confidence under 0.7	56 / 11	182 / 17	80 / 6
Retake phrases cut	4	6	5
Crew phrases cut	0	5	0
Reference words	93	220	70
Word differences from reference	0	0	1

Differences are measured on caption words. References keep the transcript's wording, so this checks which lines survive, not transcription accuracy.

TABLE 14. Learned Advantage of One Action Over Another (Posterior Means, No Prior)

Context and contrast	Snappy	Relaxed	Storyteller
"um": REMOVE − KEEP	+0.94	+0.95	+0.49
"like": REMOVE − KEEP	+0.91	−0.49	−0.89
Pause 0.6 s: REMOVE − KEEP	+0.88	−0.78	+0.38
Pause 1.5 s: REMOVE − KEEP	+1.09	+0.07	+0.48
Earlier take: CUT − KEEP	+0.63	+0.39	+0.60
Ordinary phrase: CUT − KEEP	−0.69	−1.01	−0.73

earlier. There are certainly more such problems in footage we have not tried. We also ran the editor ourselves on these recordings during development and turned what we saw into fixes; this feedback was informal and is not a user study.

9) What Did the Model Learn?

Acceptance alone does not show whether the learned weights make sense, so we queried the posterior means directly. We trained with no prior, for ten videos and six seeds, so that everything comes from feedback, and then scored hand-built contexts (Table 14). The sign of each entry says which action the model favors. The results read like the personas' hidden rules. All three creators remove "um", but only Snappy removes "like", and the others keep it. A 0.6 s pause is removed by Snappy and kept firmly by Relaxed, and Relaxed nearly flips between 0.6 s and 1.5 s, which is the threshold in its rule. An earlier take is cut and an ordinary phrase is kept by all, the behavior shared by every style. This is the basis of the profile page, and it also shows a benefit of keeping the model linear: such a check takes a few lines of code.

D. COMPARISON WITH EXISTING METHODS

Table 6 is the main benchmark. Compared with the Rule baseline, modeled on what one-click tools do, the proposed method cuts regret by 73 percent and corrections per minute of footage by 69 percent, from 19.9 to 6.2. Against the unadapted Global prior the gap is wider still, which shows that a pooled model is worse than a simple rule when styles disagree. Among learners, the differences are much smaller. Video 10 acceptance for LinUCB (0.903) and proposed (0.910) overlaps within intervals. A paired comparison on matched persona and seed puts the regret advantage of the

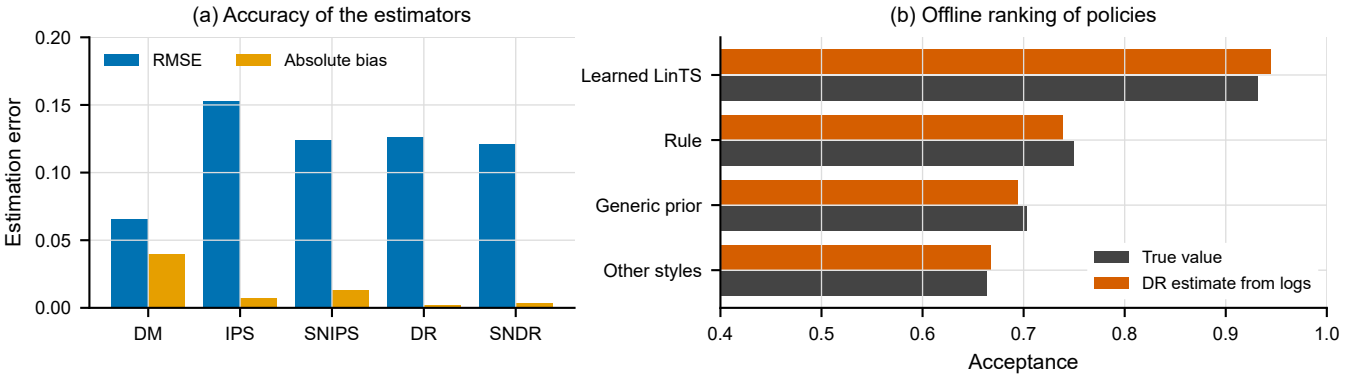


FIGURE 11. Off-policy evaluation against simulator truth. (a) Error of each estimator. (b) True value and doubly robust estimate of four target policies, computed only from the logs of the learning agent.

TABLE 15. Fifty Videos, Six Creators, Ten Seeds

Method	Video 50	Regret, 10	Regret, 50
LinTS, generic prior	0.910	194.7	765.7
Proposed, default	0.915	181.5	713.3
Proposed, retuned	0.927	**152.7**	**590.9**
Neural, blank	0.929	186.7	642.9
Neural, pretrained	0.928	167.8	618.8
NeuralLinear, blank	**0.933**	219.6	668.9
NeuralLinear, pretrained	**0.933**	202.3	659.5

proposed method at 26.6 ± 9.1 decisions over LinUCB and 13.3 ± 9.3 over plain LinTS, both with $p < 0.01$. On video 10 alone it is not distinguishable from either ($p = 0.15$ and $p = 0.22$). No correction for multiple comparisons was applied.

1) A Fair Test Against Neural Agents

Our first deep baseline made the linear model look better than it deserved. That search tried four settings, and its best ϵ was the smallest on offer, a sign that the grid was too narrow. We repeated the comparison with one procedure for every family: settings chosen on seeds 201–203 by acceptance on videos 5 and 6, and results reported on the usual seeds, with each run extended to 50 videos. The neural agent keeps one network per decision type, trained online from a replay buffer, which is the DQN recipe [2] with $\gamma = 0$ because these decisions have no transition. Its grid now has 36 settings of learning rate, ϵ, gradient steps per event, and width. We added NeuralLinear [29], which fits a Bayesian linear model on the last hidden layer and acts by Thompson sampling. Pretrained variants learn first from the other creators' labels, the same information the generic prior has. Our own method went through 60 settings of λ, v, and ϵ.

The result changed our conclusion, as Fig. 12 and Table 15 show. With wider tuning, the networks overtake our default settings: pretrained Neural makes 13.7 ± 6.9 fewer wrong proposals over ten videos and 94.5 ± 19.4 fewer over fifty. Retuning (Table 5) restores the linear method's lead. It makes 15.0 ± 6.7 fewer wrong proposals than pretrained Neural over ten videos ($p < 0.001$) and 27.9 ± 19.7 fewer over fifty ($p = 0.006$), and it starts higher, at 0.826 against 0.789. On 120 fresh runs (seeds 301–320), the retuned settings beat the defaults by 1.5 ± 0.5 points at video 10 and by 30.8 ± 3.6 wrong proposals. By video 50, however, the retuned method's acceptance is statistically level with Neural (0.927 against 0.929) and slightly below NeuralLinear (0.933, $p = 0.09$).

Retuning has a price that the acceptance figures hide. Its confidences are less useful, with calibration error 0.052 against 0.023 and an area under the curve of 0.831 against 0.893, and only 6.0 percent of suggestions fall below the review threshold instead of 10.1. A designer who relies on the review queue may reasonably keep the default, which is why most results in this paper use it. Two further cautions apply. The chosen neural learning rate was the lowest in its grid, so a still lower one might do better, and all of these creators can be described by our features, which favors a linear model. The case for the linear method therefore rests on fewer early mistakes, cost, exact replay, and readable explanations. It is not a claim that networks are worse, and with enough feedback a network is a reasonable choice.

2) Visual Rhythm: Is a Sequential Agent Needed?

Table 16 and Fig. 13 test the sequential agent on four simulated creators whose preferences depend on history: one modest punch-in after a quiet stretch, a strong beat only on important phrases after eight quiet seconds, a beat every three seconds without repeating the kind, and almost nothing. Each edits eight videos with feedback on every beat and rates each whole video by how close its beats per minute are to what the creator likes. Scores are on four unseen videos, over 20 seeds not used for tuning. DQN and PPO were tuned over learning rate and ϵ or update epochs on separate seeds, from a blank start and from a start fitted to the same rule that gives RLSVI its prior.

RLSVI did not beat the bandit. Its final match is 0.901 against 0.919 (paired difference −0.018 ± 0.019, $p = 0.055$), and their ratings are level at +0.04. We expected the sequential agent to win, and the likely reason it did not lies in

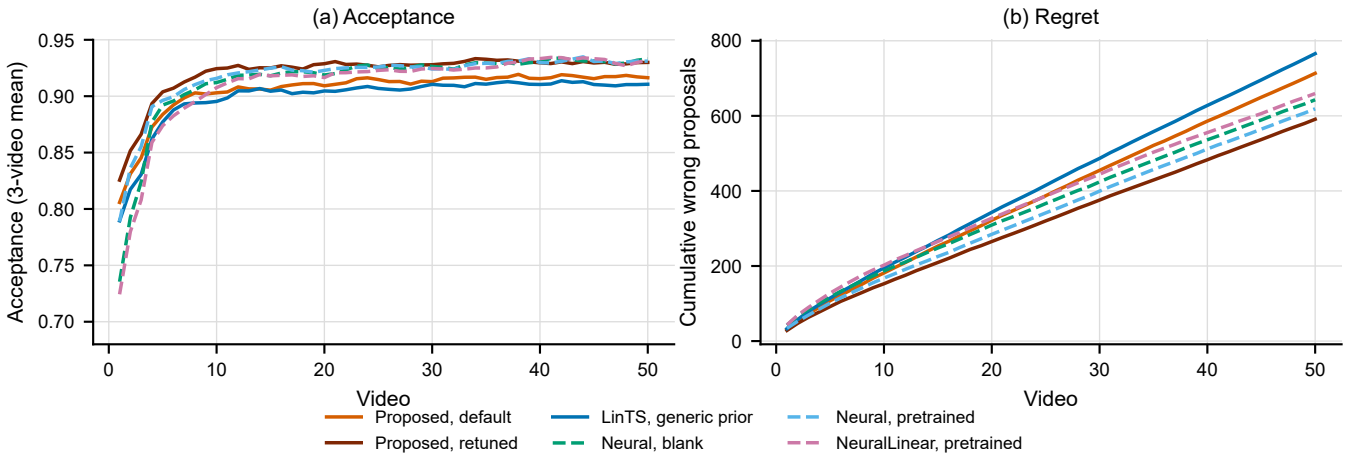


FIGURE 12. Comparison over 50 videos with every family tuned on the same separate seeds. The retuned method leads early and keeps the lowest cumulative regret, while network-based agents draw level on acceptance.

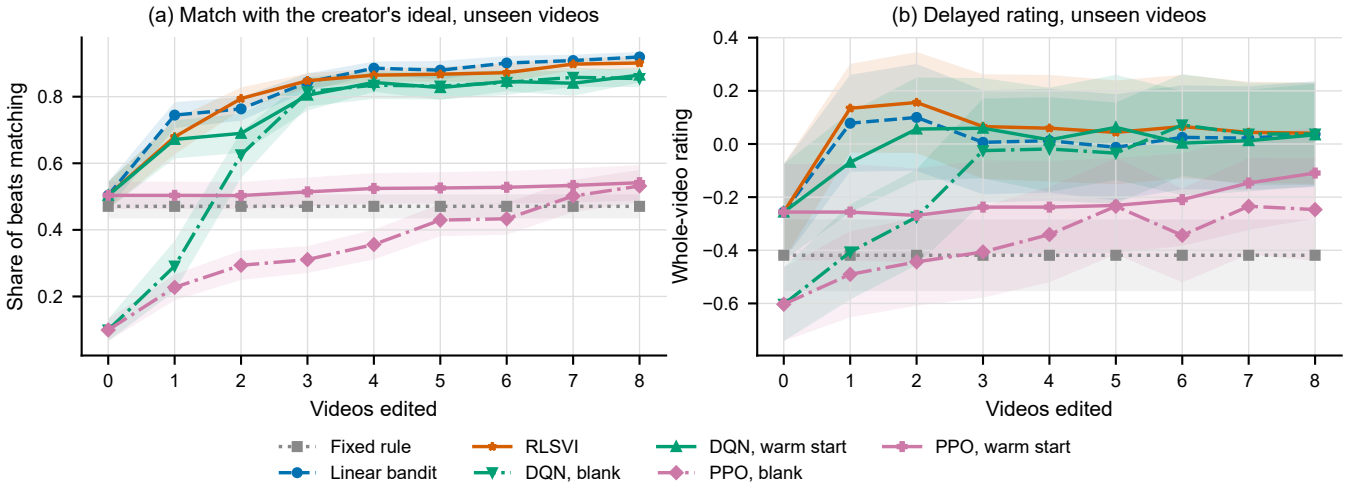


FIGURE 13. Visual rhythm on unseen videos. (a) Share of beats that match the creator's ideal. (b) Whole-video rating, in $[-1, 1]$. Bands show 95 percent intervals over 80 runs per agent.

TABLE 16. Visual Rhythm on Unseen Videos (Four Creators, 20 Seeds)

Agent	Match, video 3	Match, video 8	Rating, video 8
Fixed rule	0.471	0.471 ± 0.036	-0.42 ± 0.13
Linear bandit	0.843	0.919 ± 0.015	+0.04 ± 0.20
RLSVI	0.848	0.901 ± 0.020	+0.04 ± 0.19
DQN, blank	0.815	0.855 ± 0.026	+0.04 ± 0.19
DQN, warm start	0.805	0.865 ± 0.023	+0.03 ± 0.19
PPO, blank	0.311	0.532 ± 0.048	-0.25 ± 0.19
PPO, warm start	0.514	0.542 ± 0.053	-0.11 ± 0.17

our design. Every state these creators react to is also visible to the bandit through the history features, so choosing the best action for the current state is already close to optimal. The delayed rating adds little that per-beat feedback has not already said. The largest difference is on the minimal creator (0.816 against 0.894), where the agent chases rare positive beats. Both agents still fail on rating for the two creators who want few beats, because a few wrong beats are enough to miss a narrow target rate. The deep agents had far less to work with, about 260 samples per run. The better DQN reached 0.865 and the better PPO 0.542, so at this data budget both trail the linear agents on matches even after tuning, although DQN's final rating is level with theirs and every interval on rating is wide. These creators give feedback on every beat, which leaves little for a delayed reward to add.

To give RLSVI its fair test, we reran the agents for 16 videos while withholding per-beat feedback, down to the whole-video rating alone (Table 17). Here the picture splits. RLSVI's rating overtakes the bandit's as feedback thins, by 0.21 ± 0.11 at 10 percent and by 0.48 ± 0.15 at rating only (both $p < 0.001$), while its beat match falls well behind (0.336 against 0.515). In other words, it learns how often a creator wants beats without learning where to put them. Part of this rating gap is built in, because a bandit in this setup cannot use the rating at all. DQN receives the rating too and

TABLE 17. Visual Rhythm With Sparse Feedback (Video 16, Match / Rating)

Beats with feedback	RLSVI	Bandit	DQN, warm
100%	0.915 / +0.06	0.945 / +0.03	0.906 / +0.00
30%	0.825 / +0.18	0.910 / +0.06	0.808 / +0.04
10%	0.687 / +0.24	0.828 / +0.04	0.670 / +0.01
Rating only	0.336 / +0.16	0.515 / −0.32	0.276 / −0.31

TABLE 18. Ablation: Change in Acceptance Against the Full Method

Removed	Videos 1–3	Videos 8–10	Regret
Nothing (full method)	0.845	0.903	181.5
Manual-edit mapping	−0.025	−0.011	206.5
Prior altogether	−0.021	−0.003	198.2
Retake features	−0.021	−0.016	213.1
Style prior (generic kept)	−0.015	−0.008	194.7
Semantic features	−0.003	−0.003	190.3
Implicit feedback	+0.004	+0.002	177.9
History features	+0.005	−0.005	186.2
Discounting	+0.003	+0.018	168.5
Exploration	+0.027	+0.026	140.8

does not benefit (rating −0.31). All settings were tuned under full feedback, so these figures may understate both deep and randomized agents.

E. ABLATION STUDY

Table 18 and Fig. 14 remove one component at a time, with paired seeds. Three components make clear positive contributions. Manual-edit mapping adds 2.5 points on early videos and 1.1 late, as half of the corrections would otherwise be lost. Retake features add 2.1 points early and 1.6 late, the largest lasting effect, as expected since every persona cuts earlier takes. The style prior adds 1.5 points early and 0.8 late. Removing the prior altogether gives the largest early drop among prior variants (2.1) with no significant late penalty. Semantic and history features have almost no effect in the simulator. For semantic features that may be partly an artifact, because synthetic embeddings are only crude stand-ins for language.

Three removals do not hurt, and two of them improve the score, which runs against the usual narrative. Removing implicit feedback changes little (+0.2 late, within noise). Removing discounting adds 1.8 points late, which is expected since these personas do not drift, and Fig. 8 shows the opposing case. The more important result is exploration: a policy with $v = 0$ and $\epsilon = 0$ scores 2.7 points higher early and 2.6 later, with regret of 140.8 against 181.5.

Acceptance during training counts every exploratory proposal as a mistake, so it cannot say whether exploration led to a better policy. We therefore trained each variant on ten videos under three feedback conditions, froze its posterior mean, and scored it on four unseen videos with no corrections. Table 19 gives the outcome over 48 runs per cell. Held-out acceptance barely depends on the exploration strategy: every paired difference from Thompson sampling at $v = 0.25$ lies within about one point, and most intervals include zero. A purely greedy learner ends at 0.931 against 0.932 under normal feedback. The training cost, on the other

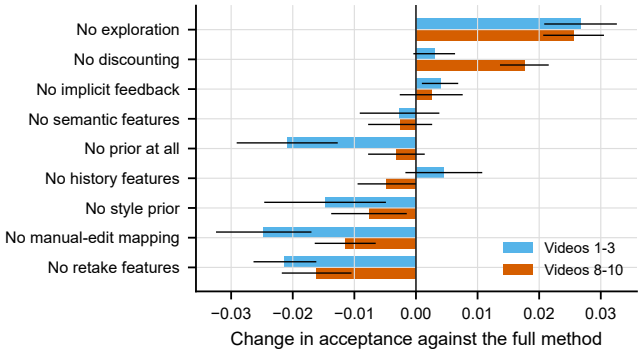


FIGURE 14. Paired change in acceptance from removing each component. Negative bars mean that the component helps. Error bars are 95 percent intervals.

TABLE 19. Exploration Variants: Acceptance During Training and Frozen on Unseen Videos

	Normal		Reluctant		Review 30%	
Variant	Train	Held	Train	Held	Train	Held
Greedy	0.908	0.931	0.884	0.917	0.867	0.908
Thompson, $v = 0.05$	0.897	0.934	0.873	0.914	0.852	0.902
Thompson, $v = 0.25$	0.883	0.932	0.855	0.921	0.837	0.900
Thompson, $v = 0.5$	0.851	0.933	0.810	0.918	0.800	0.906
LinUCB	0.880	0.936	0.841	0.921	0.825	0.911
ϵ-greedy, $\epsilon = 0.1$	0.860	0.932	0.832	0.913	0.816	0.903

Reluctant: correction probability 0.2. Held-out intervals are about ±0.01 throughout.

hand, is large and visible to the creator, with 0.908 for greedy against 0.883 for $v = 0.25$ and 0.851 for $v = 0.5$.

Why does the usual argument for exploration not apply here? Three reasons stand out. Each video presents fresh contexts, so even a greedy agent sees many feature combinations. A creator's override names the correct arm and updates both arms, so the data are close to full information. And a prior mean already provides an informed start. The simulator does not contain rejections that withhold the right answer, which is the setting where exploration should matter most, so we can only say that exploration has not been shown to help under the feedback we modeled. In the deployed configuration Thompson sampling is kept for two practical reasons: it provides the confidence that orders the review queue, and its stochastic choices give the logged propensities that off-policy evaluation needs. A greedy logger would leave nothing to reweight. A smaller v than ours scores higher on acceptance, but Table 11 shows what it gives up.

F. SENSITIVITY ANALYSIS

Fig. 15 varies one hyperparameter at a time over six seeds per persona. Prior strength λ has a broad optimum near 2 to 10 late in training. Weak priors ($\lambda = 0.5$) lose the cold-start benefit, at 0.818 early, and very strong ones ($\lambda = 30$) learn too slowly (0.807 early, 0.884 late). That loss is what the product accepts for keep/cut alone, to prevent one click from causing large changes. Posterior scale v has the largest effect of any parameter: at $v = 2$ the agent is little better than a rule (0.697 late), and it improves monotonically down to the

smallest value tried ($v = 0.05$, 0.916 late). The exploration floor ϵ behaves in the same way, from 0.913 at 0 to 0.819 at 0.2. Discount γ gives 0.918 late at 1.0 and 0.858 at 0.95. The implicit reward has a small effect, with lower values marginally better (0.909 at 0, 0.893 at 1).

The effect of seeds, the last factor, is modest. Late acceptance intervals are about ± 0.012 throughout, so differences below two points between settings should not be read as meaningful. In all five panels, what improves stationary acceptance is less randomness and less forgetting, which is the same message as the ablation.

G. ROBUSTNESS AND GENERALIZATION

Robustness asks what happens when the environment gets worse. Fig. 16 shows three such stresses. With the label noise raised from 0 to 30 percent, acceptance on videos 8 to 10 falls steadily from 0.920 to 0.749, staying well above Rule (about 0.59). A creator who reviews just 10 percent of decisions still leads the agent to 0.825, and 30 percent reaches 0.873, against 0.902 at full review. Reluctance to correct matters little: dropping the correction probability from 1.0 to 0.2 lowers acceptance from 0.910 to 0.889, because accepts still carry information. Video length also has little effect. With 12, 30, and 60 phrases per video (about 67, 160, and 315 decisions), acceptance is 0.898, 0.902, and 0.912, consistent with the idea that learning is measured in decisions rather than videos.

Generalization asks about conditions the model and its prior were not built for. Fig. 17 tests creators who were deliberately placed outside the model family and outside the training data of the prior, scored only on keep/cut, pause, and filler decisions since these personas never zoom. By video 8 (Fig. 17), LinTS reaches 0.883 with nonlinear rules, well above Rule (0.511) and the best fixed choice per type (0.489), because binned features can approximate interval rules. With noisy feedback (20 percent wrong labels, 30 percent correction rate) it reaches 0.810, still below the best fixed arm at 0.835. With arbitrary rules that depend on a hash of the unit, it ends at 0.499, below the fixed arm at 0.568, because nothing can be learned there and the agent chases noise.

Table 20 extends these creators to 30 videos and adds the tuned networks, with eight seeds per cell, so intervals are wide. Each family wins where its assumptions hold. On nonlinear rules the network makes roughly half as many mistakes as LinTS. Under noise the order reverses: LinTS passes the fixed arm and beats blank Neural by 7.3 ± 3.8 points. On arbitrary rules every learner stays below a constant policy, so the method is not winning by construction.

H. REPRODUCIBILITY

Every simulated result can be regenerated from the repository, which contains the simulator, the personas, and the experiment scripts behind each table and figure. A complete rerun of the comparison, ablation, sensitivity, and robustness groups involves about 2500 simulated creators and takes under half an hour on the laptop described above, and the final deep comparison takes about an hour and a half. Seeds are fixed per run. Synthetic videos are deterministic functions of their seed, and all methods compared at a given seed see the same videos, which is what makes the paired comparisons valid. Priors are fitted on separate videos and with the evaluated persona excluded. The real clips cannot be redistributed without the rights holder's consent, so the real-audio results can be repeated only with other footage. Repeating the analysis three times on the two real clips in fresh processes gave identical transcripts and cuts. All reported intervals are 95 percent t intervals across runs, and we state explicitly where a comparison is paired.

TABLE 20. Creators Outside the Model Family, 30 Videos (Acceptance / Wrong Proposals)

Creator	LinTS	Neural	Neural, pretr.	Fixed
Nonlinear	0.931 / 454	0.948 / 244	0.953 / 230	0.479
Noisy	0.953 / 541	0.880 / 564	0.873 / 532	0.836
Arbitrary	0.508 / 1159	0.511 / 1140	0.477 / 1142	0.559

Acceptance counts keep/cut, pause, and filler decisions; wrong proposals count all decisions.

I. HOW MUCH OF THIS IS REINFORCEMENT LEARNING?

A fair reader may ask whether this is RL or online regression with a user interface, so we address the question directly. Several properties belong to RL and not to supervised learning: the agent acts and observes feedback only on what it proposed, its data depend on its own past actions, it holds an explicit posterior and chooses by sampling from it, it learns online per user under non-stationary preferences, it decides what to ask about, it logs propensities that make counterfactual evaluation possible, and it plans over an episode with a learned reward, and in the rhythm agent its actions change the state it then acts in. A contextual bandit is the one-state special case of an RL problem [1], [32], and our planner is an MDP solved by posterior sampling.

Other properties weaken the claim. Keep/cut and filler decisions are binary, so a rejection reveals the right label, and in those cases the problem is close to online supervised learning. Our own experiments show that exploration does not improve held-out policies under this feedback. In most decisions actions affect later states only weakly, through history features that also proved to have little effect, and where the state does matter, a sequential agent beat a bandit only on whole-video rating under scarce feedback. The method is therefore best described as a contextual bandit system with a planning layer, and not as deep or fully sequential RL. We consider that the right tool for the amount of data available, and we think it is more useful to state its boundaries than to use a larger model that could not be justified by a few hundred human signals.

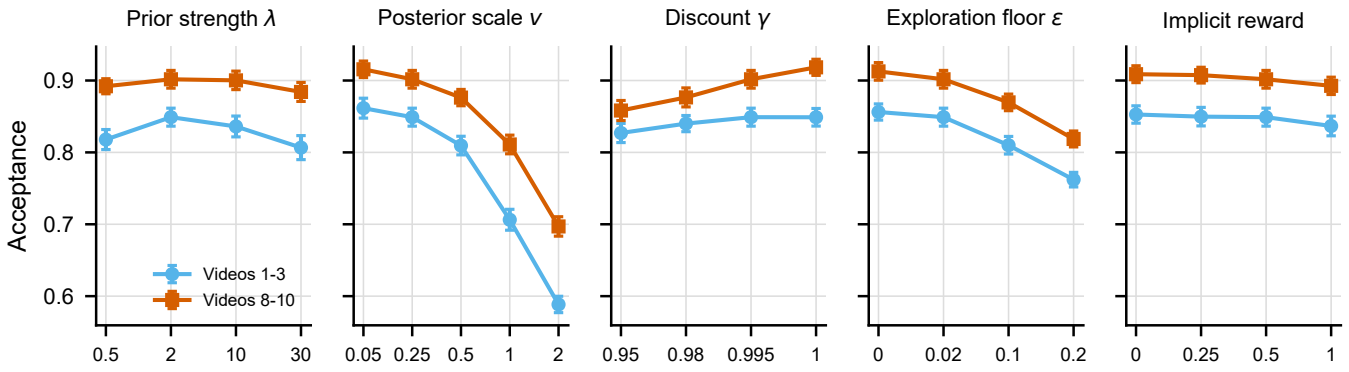


FIGURE 15. One-at-a-time sensitivity of acceptance on early and late videos. Posterior scale and exploration floor have the largest effects; the others are comparatively flat near their chosen values.

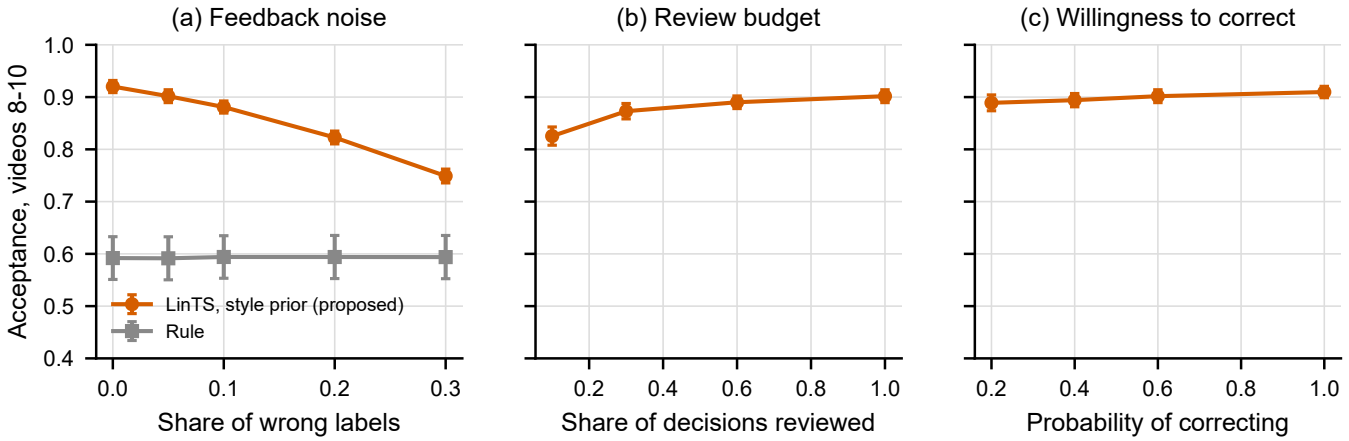


FIGURE 16. Acceptance on videos 8 to 10 under (a) wrong feedback labels, (b) smaller review budgets, and (c) reluctance to correct. Degradation is gradual in all three cases.

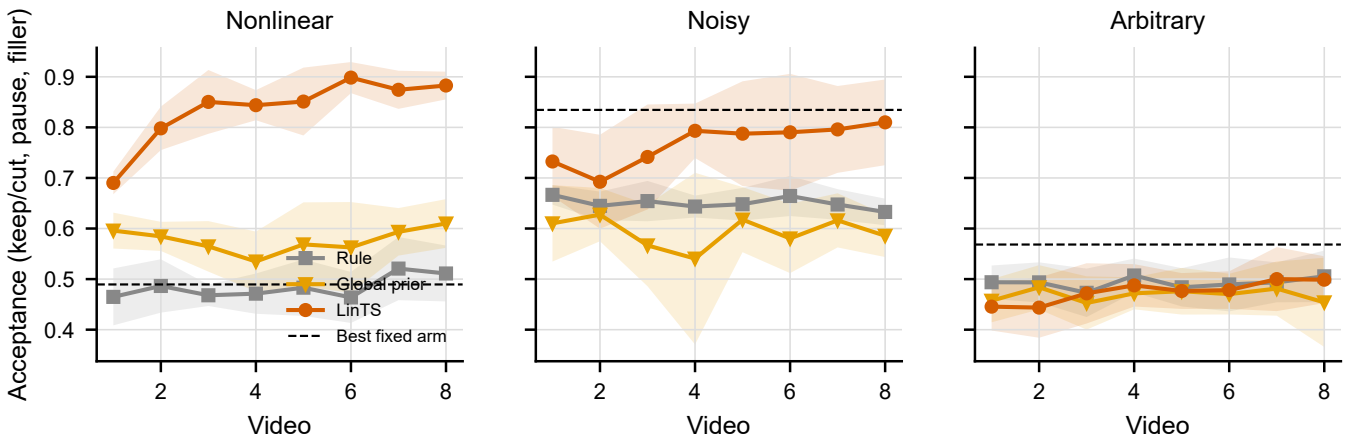


FIGURE 17. Creators outside the model family and the prior's training data. The first eight videos. The method generalizes to nonlinear preferences but is still behind a constant policy under heavy noise and arbitrary preferences.

J. THREATS TO VALIDITY

Three threats deserve separate mention. Internal validity is threatened by the fact that the experiments share one code base with the product and were run while that code was still changing. We reran the simulator experiments in one session for this paper, and the scripts in the repository regenerate every table and figure, but the analysis code was still being revised for real audio during that period. Hyperparameters were tuned on the same simulator used for evaluation, which favors the method. For the later results we separated this:

rhythm settings, the deep baselines, and the retuned linear settings were tuned on seeds not used for reporting, and the retuned and product settings were confirmed on fresh seeds. The default $v = 0.25$ was selected for the quality of its confidences and not for peak acceptance, so it does not inflate the headline figure. We cannot rule out that earlier tuning affected other results.

Construct validity concerns whether acceptance is the right target. A creator might accept a mediocre suggestion simply because correcting takes effort, and in that case acceptance would overstate quality. Our reward distinguishes an explicit accept from an untouched suggestion for this reason, but it cannot see what a person would have preferred had they thought harder. Equally, matching a creator's past decisions is not the same as making a video that audiences enjoy. We optimize the former, as we have no access to the latter.

External validity is the largest threat. Synthetic transcripts contain the mistakes we thought to plant, in the proportions we chose. Their semantic embeddings are random vectors around a topic and not real sentence embeddings, so the near-zero effect of semantic features in the ablation may not carry over to real speech. The result for arbitrary preferences shows that the method can lose when its assumptions fail, and a real population will lie somewhere between our personas and that extreme.

IV. LIMITATIONS

The most serious limitation is that no formal human study has been run. Every learning result concerns simulated creators that we designed, using features that we also designed, and our personas' rules are largely expressible in those features. The tests with creators outside the model family push against this, but the acceptance numbers should not be read as predictions for people. Real creators may be less consistent, may tire of reviewing, and may hold preferences that are not functions of these features. The study comparing time-to-post against a simple editor, which is the real test of the product claim, is still future work. Real logs in our database are also too few for off-policy conclusions about real use.

The analysis has been checked on only three clips, two real and one with a synthetic voice, against references prepared from the transcripts, with word-level comparison and no formal listening test. The crew rules were written by looking at one of these clips. Only English is supported; other languages get a warning, not a usable edit. Running Whisper in a fresh process removes a difference we could reproduce but not explain, so it is a workaround and not a diagnosis. Fillers with no audio dip stay in the edit. Disfluency recovery can miss repeats that the CTC model does not read clearly, and Whisper can still mishear words such as names.

On the RL side, binary keep/cut and filler decisions are close to online supervised learning, particularly because a correction reveals the right label. The bandit layer treats decisions as nearly independent, and the planner's frozen history features remove most of the sequential interaction from its reward. So the planner's sequential content is a length budget, and the rhythm agent, the richer sequential model, has only been tested against simulated creators whose preferences a bandit can also see. The planner also does not improve which phrases are chosen. Off-policy estimates are limited by an effective sample size of 15 percent and intervals that under-cover, and the removal of the exploration floor on keep/cut in the product means that actions Thompson sampling never picks have no support in real logs.

Scalability is not a concern for the learner, but analysis takes one to two times the clip's length on a laptop CPU, and uploads are limited to five minutes. There is also a safety aspect, in that a wrong CUT silently deletes content on an autopilot export. Autopilot therefore shows a badge based on expected corrections and a trust score, but we have not validated that badge.

V. CONCLUSION AND FUTURE WORK

This paper treated talking-head editing as online RL from one creator's normal editing behavior, a setting in which rule-based tools cannot adapt and supervised models lack data. We combined a linear Thompson sampling bandit with a tempered style posterior for the cold start, an uncertainty-ordered review queue, an exact posterior-sampling planner for length targets, a randomized value-iteration agent for visual pacing, and propensity logging for offline evaluation. In simulation, retuned settings reach 93.1 percent acceptance by the tenth video against 58.6 percent for a rule baseline, and they make the fewest wrong proposals over 50 videos of any agent we tried. Learned policies are strongly creator-specific (0.936 against 0.468 across creators), review at a 30 percent budget gives 86.9 percent correct decisions against 63.3 percent in timeline order, and the planner reduces length error from 40.0 to 4.6 percent for 0.4 points of acceptance.

The experiments also set limits on these claims. Fairly tuned networks overtook our default settings and drew level with the retuned ones on acceptance by the fiftieth video, and they did better on nonlinear taste, so the linear choice rests on scarce-feedback performance, cost, and transparency. Retuning for fewer mistakes weakened the confidence signal that the review queue needs. Random exploration did not prove valuable here, because corrections carry enough information for a nearly greedy policy to do well, and value iteration improved on a bandit only for whole-video rating under scarce feedback. Nothing is yet established about real creators.

Future work follows directly. A study with real creators, comparing time-to-post and blind quality ratings against a simple editor, comes first, along with off-policy evaluation on those logs. A larger set of real clips is needed with independent reference edits and listening tests. The keep/cut selection under a length target needs a better reward, possibly through pairwise preference feedback [7]. History features could return to the planner with uncertainty penalties [45], now that we know how they fail, and RLSVI should be tuned for sparse feedback so that its rating gain can be tested with people. Crew-chatter detection needs a labeled set and a learned model. Exploration should be restricted to where

it is useful, such as multi-arm decisions with rejection-only feedback. Style prototypes could eventually be learned from real users' posteriors rather than written by hand, as in meta-learning approaches [34], [35].

ACKNOWLEDGMENT

The authors thank Prof. Geetha M for her guidance. Claude, an AI model from Anthropic, was used to assist with software development, experiments, and drafting and editing this manuscript. The authors reviewed the content and take responsibility for it.

REFERENCES

[1] R. S. Sutton and A. G. Barto, Reinforcement Learning: An Introduction, 2nd ed. Cambridge, MA, USA: MIT Press, 2018.
[2] V. Mnih et al., "Human-level control through deep reinforcement learning," Nature, vol. 518, no. 7540, pp. 529–533, Feb. 2015, doi: 10.1038/nature14236.
[3] L. Li, W. Chu, J. Langford, and R. E. Schapire, "A contextual-bandit approach to personalized news article recommendation," in Proc. 19th Int. Conf. World Wide Web, Raleigh, NC, USA, 2010, pp. 661–670, doi: 10.1145/1772690.1772758.
[4] G. Zheng et al., "DRN: A deep reinforcement learning framework for news recommendation," in Proc. World Wide Web Conf., Lyon, France, 2018, pp. 167–176, doi: 10.1145/3178876.3185994.
[5] M. Chen, A. Beutel, P. Covington, S. Jain, F. Belletti, and E. H. Chi, "Top-K off-policy correction for a REINFORCE recommender system," in Proc. 12th ACM Int. Conf. Web Search Data Mining, Melbourne, VIC, Australia, 2019, pp. 456–464, doi: 10.1145/3289600.3290999.
[6] Y. Hu, H. He, C. Xu, B. Wang, and S. Lin, "Exposure: A white-box photo post-processing framework," ACM Trans. Graph., vol. 37, no. 2, pp. 1–17, Apr. 2018, doi: 10.1145/3181974.
[7] P. F. Christiano, J. Leike, T. B. Brown, M. Martic, S. Legg, and D. Amodei, "Deep reinforcement learning from human preferences," in Proc. Adv. Neural Inf. Process. Syst., vol. 30, 2017, pp. 4299–4307.
[8] L. Ouyang et al., "Training language models to follow instructions with human feedback," in Proc. Adv. Neural Inf. Process. Syst., vol. 35, 2022, pp. 27730–27744.
[9] F. Berthouzoz, W. Li, and M. Agrawala, "Tools for placing cuts and transitions in interview video," ACM Trans. Graph., vol. 31, no. 4, pp. 1–8, Jul. 2012, doi: 10.1145/2185520.2185563.
[10] S. Rubin, F. Berthouzoz, G. J. Mysore, W. Li, and M. Agrawala, "Content-based tools for editing audio stories," in Proc. 26th ACM Symp. User Interface Softw. Technol., St. Andrews, U.K., 2013, pp. 113–122, doi: 10.1145/2501988.2501993.
[11] B. Huber, H. V. Shin, B. Russell, O. Wang, and G. J. Mysore, "B-Script: Transcript-based B-roll video editing with recommendations," in Proc. CHI Conf. Hum. Factors Comput. Syst., Glasgow, U.K., 2019, pp. 1–11, doi: 10.1145/3290605.3300311.
[12] W. Chu, L. Li, L. Reyzin, and R. E. Schapire, "Contextual bandits with linear payoff functions," in Proc. 14th Int. Conf. Artif. Intell. Statist., Fort Lauderdale, FL, USA, 2011, pp. 208–214.
[13] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," 2017, arXiv:1707.06347.
[14] H. V. Shin, W. Li, and F. Durand, "Dynamic authoring of audio with linked scripts," in Proc. 29th Annu. Symp. User Interface Softw. Technol., Tokyo, Japan, 2016, pp. 509–516, doi: 10.1145/2984511.2984561.
[15] A. Truong, F. Berthouzoz, W. Li, and M. Agrawala, "QuickCut: An interactive tool for editing narrated video," in Proc. 29th Annu. Symp. User Interface Softw. Technol., Tokyo, Japan, 2016, pp. 497–507, doi: 10.1145/2984511.2984569.
[16] A. Pavel, C. Reed, B. Hartmann, and M. Agrawala, "Video digests: A browsable, skimmable format for informational lecture videos," in Proc. 27th Annu. ACM Symp. User Interface Softw. Technol., Honolulu, HI, USA, 2014, pp. 573–582, doi: 10.1145/2642918.2647400.
[17] M. Leake, A. Davis, A. Truong, and M. Agrawala, "Computational video editing for dialogue-driven scenes," ACM Trans. Graph., vol. 36, no. 4, pp. 1–14, Jul. 2017, doi: 10.1145/3072959.3073653.
[18] O. Fried et al., "Text-based editing of talking-head video," ACM Trans. Graph., vol. 38, no. 4, pp. 1–14, Jul. 2019, doi: 10.1145/3306346.3323028.
[19] B. Wang, Y. Li, Z. Lv, H. Xia, Y. Xu, and R. Sodhi, "LAVE: LLM-powered agent assistance and language augmentation for video editing," in Proc. 29th Int. Conf. Intell. User Interfaces, Greenville, SC, USA, 2024, pp. 699–714, doi: 10.1145/3640543.3645143.
[20] M. Huh, D. Li, K. Pimmel, H. V. Shin, A. Pavel, and M. Dontcheva, "VideoDiff: Human-AI video co-creation with alternatives," in Proc. CHI Conf. Hum. Factors Comput. Syst., Yokohama, Japan, 2025, pp. 1–19, doi: 10.1145/3706598.3713417.
[21] M. Wang, G.-W. Yang, S.-M. Hu, S.-T. Yau, and A. Shamir, "Write-A-Video: Computational video montage from themed text," ACM Trans. Graph., vol. 38, no. 6, pp. 1–13, Nov. 2019, doi: 10.1145/3355089.3356520.
[22] K. Zhou, Y. Qiao, and T. Xiang, "Deep reinforcement learning for unsupervised video summarization with diversity-representativeness reward," in Proc. AAAI Conf. Artif. Intell., vol. 32, no. 1, 2018, doi: 10.1609/aaai.v32i1.12255.
[23] P. Hu, N. Xiao, F. Li, Y. Chen, and R. Huang, "A reinforcement learning-based automatic video editing method using pre-trained vision-language model," in Proc. 31st ACM Int. Conf. Multimedia, Ottawa, ON, Canada, 2023, pp. 6441–6450, doi: 10.1145/3581783.3611878.
[24] W. R. Thompson, "On the likelihood that one unknown probability exceeds another in view of the evidence of two samples," Biometrika, vol. 25, nos. 3–4, pp. 285–294, Dec. 1933, doi: 10.2307/2332286.
[25] E. Kaufmann, N. Korda, and R. Munos, "Thompson sampling: An asymptotically optimal finite-time analysis," in Proc. Int. Conf. Algorithmic Learn. Theory, Lyon, France, 2012, pp. 199–213, doi: 10.1007/978-3-642-34106-9_18.
[26] S. Agrawal and N. Goyal, "Thompson sampling for contextual bandits with linear payoffs," in Proc. 30th Int. Conf. Mach. Learn., Atlanta, GA, USA, 2013, pp. 127–135.
[27] M. Abeille and A. Lazaric, "Linear Thompson sampling revisited," Electron. J. Statist., vol. 11, no. 2, pp. 5165–5197, 2017, doi: 10.1214/17-EJS1341SI.
[28] D. J. Russo, B. Van Roy, A. Kazerouni, I. Osband, and Z. Wen, "A tutorial on Thompson sampling," Found. Trends Mach. Learn., vol. 11, no. 1, pp. 1–96, 2018, doi: 10.1561/2200000070.
[29] C. Riquelme, G. Tucker, and J. Snoek, "Deep Bayesian bandits showdown: An empirical comparison of Bayesian deep networks for Thompson sampling," in Proc. Int. Conf. Learn. Represent., Vancouver, BC, Canada, 2018.
[30] D. Zhou, L. Li, and Q. Gu, "Neural contextual bandits with UCB-based exploration," in Proc. 37th Int. Conf. Mach. Learn., 2020, pp. 11492–11502.
[31] D. Bouneffouf, I. Rish, and C. Aggarwal, "Survey on applications of multi-armed and contextual bandits," in Proc. IEEE Congr. Evol. Comput., Glasgow, U.K., 2020, pp. 1–8, doi: 10.1109/CEC48606.2020.9185782.
[32] T. Lattimore and C. Szepesvári, Bandit Algorithms. Cambridge, U.K.: Cambridge Univ. Press, 2020, doi: 10.1017/9781108571401.
[33] Y. Russac, C. Vernade, and O. Cappé, "Weighted linear bandits for non-stationary environments," in Proc. Adv. Neural Inf. Process. Syst., vol. 32, 2019.
[34] B. Kveton et al., "Meta-Thompson sampling," in Proc. 38th Int. Conf. Mach. Learn., 2021, pp. 5884–5893.
[35] L. Cella, A. Lazaric, and M. Pontil, "Meta-learning with stochastic linear bandits," in Proc. 37th Int. Conf. Mach. Learn., 2020, pp. 1360–1370.
[36] J. Hong, B. Kveton, M. Zaheer, Y. Chow, A. Ahmed, and C. Boutilier, "Latent bandits revisited," in Proc. Adv. Neural Inf. Process. Syst., vol. 33, 2020, pp. 13423–13433.
[37] W. B. Knox and P. Stone, "Interactively shaping agents via human reinforcement: The TAMER framework," in Proc. 5th Int. Conf. Knowl. Capture, Redondo Beach, CA, USA, 2009, pp. 9–16, doi: 10.1145/1597735.1597738.
[38] J. MacGlashan et al., "Interactive learning from policy-dependent human feedback," in Proc. 34th Int. Conf. Mach. Learn., Sydney, NSW, Australia, 2017, pp. 2285–2294.
[39] Y. Hu, Y. Koren, and C. Volinsky, "Collaborative filtering for implicit feedback datasets," in Proc. 8th IEEE Int. Conf. Data Mining, Pisa, Italy, 2008, pp. 263–272, doi: 10.1109/ICDM.2008.22.
[40] B. Settles, Active Learning (Synthesis Lectures on Artificial Intelligence and Machine Learning). Cham, Switzerland: Springer, 2012, doi: 10.1007/978-3-031-01560-1.
[41] I. Osband, D. Russo, and B. Van Roy, "(More) efficient reinforcement learning via posterior sampling," in Proc. Adv. Neural Inf. Process. Syst., vol. 26, 2013, pp. 3003–3011.

[42] I. Osband, B. Van Roy, and Z. Wen, "Generalization and exploration via randomized value functions," in Proc. 33rd Int. Conf. Mach. Learn., New York, NY, USA, 2016, pp. 2377–2386.
[43] I. Osband, B. Van Roy, D. J. Russo, and Z. Wen, "Deep exploration via randomized value functions," J. Mach. Learn. Res., vol. 20, no. 124, pp. 1–62, 2019.
[44] M. Janner, J. Fu, M. Zhang, and S. Levine, "When to trust your model: Model-based policy optimization," in Proc. Adv. Neural Inf. Process. Syst., vol. 32, 2019, pp. 12498–12509.
[45] T. Yu et al., "MOPO: Model-based offline policy optimization," in Proc. Adv. Neural Inf. Process. Syst., vol. 33, 2020, pp. 14129–14142.
[46] E. Altman, Constrained Markov Decision Processes. Boca Raton, FL, USA: Chapman & Hall/CRC, 1999.
[47] L. Li, W. Chu, J. Langford, and X. Wang, "Unbiased offline evaluation of contextual-bandit-based news article recommendation algorithms," in Proc. 4th ACM Int. Conf. Web Search Data Mining, Hong Kong, 2011, pp. 297–306, doi: 10.1145/1935826.1935878.
[48] Y. Saito and T. Joachims, "Counterfactual learning and evaluation for recommender systems: Foundations, implementations, and recent advances," in Proc. 15th ACM Conf. Recommender Syst., Amsterdam, The Netherlands, 2021, pp. 828–830, doi: 10.1145/3460231.3473320.
[49] D. G. Horvitz and D. J. Thompson, "A generalization of sampling without replacement from a finite universe," J. Amer. Statist. Assoc., vol. 47, no. 260, pp. 663–685, Dec. 1952, doi: 10.1080/01621459.1952.10483446.
[50] A. Swaminathan and T. Joachims, "The self-normalized estimator for counterfactual learning," in Proc. Adv. Neural Inf. Process. Syst., vol. 28, 2015, pp. 3231–3239.
[51] M. Dudík, J. Langford, and L. Li, "Doubly robust policy evaluation and learning," in Proc. 28th Int. Conf. Mach. Learn., Bellevue, WA, USA, 2011, pp. 1097–1104.
[52] A. Swaminathan et al., "Off-policy evaluation for slate recommendation," in Proc. Adv. Neural Inf. Process. Syst., vol. 30, 2017, pp. 3632–3642.
[53] V. Hadad, D. A. Hirshberg, R. Zhan, S. Wager, and S. Athey, "Confidence intervals for policy evaluation in adaptive experiments," Proc. Nat. Acad. Sci. USA, vol. 118, no. 15, 2021, Art. no. e2014602118, doi: 10.1073/pnas.2014602118.
[54] A. C. Cameron, J. B. Gelbach, and D. L. Miller, "Bootstrap-based improvements for inference with clustered errors," Rev. Econ. Statist., vol. 90, no. 3, pp. 414–427, Aug. 2008, doi: 10.1162/rest.90.3.414.
[55] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, "Robust speech recognition via large-scale weak supervision," in Proc. 40th Int. Conf. Mach. Learn., Honolulu, HI, USA, 2023, pp. 28492–28518.
[56] G. Smith, D. Yee, J. K. Chen, and L. Findlater, "Prompting Whisper for improved verbatim transcription and end-to-end miscue detection," in Proc. Interspeech, Rotterdam, The Netherlands, 2025, pp. 1943–1947, doi: 10.21437/Interspeech.2025-1031.
[57] A. Baevski, H. Zhou, A. Mohamed, and M. Auli, "wav2vec 2.0: A framework for self-supervised learning of speech representations," in Proc. Adv. Neural Inf. Process. Syst., vol. 33, 2020, pp. 12449–12460.
[58] A. Graves, S. Fernández, F. Gomez, and J. Schmidhuber, "Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks," in Proc. 23rd Int. Conf. Mach. Learn., Pittsburgh, PA, USA, 2006, pp. 369–376, doi: 10.1145/1143844.1143891.
[59] M. Bain, J. Huh, T. Han, and A. Zisserman, "WhisperX: Time-accurate speech transcription of long-form audio," in Proc. Interspeech, Dublin, Ireland, 2023, pp. 4489–4493, doi: 10.21437/Interspeech.2023-78.
[60] V. Zayats, M. Ostendorf, and H. Hajishirzi, "Disfluency detection using a bidirectional LSTM," in Proc. Interspeech, San Francisco, CA, USA, 2016, pp. 2523–2527, doi: 10.21437/Interspeech.2016-1247.
[61] M. Honnibal and M. Johnson, "Joint incremental disfluency detection and dependency parsing," Trans. Assoc. Comput. Linguistics, vol. 2, pp. 131–142, 2014, doi: 10.1162/tacl_a_00171.
[62] P. J. Lou and M. Johnson, "End-to-end speech recognition and disfluency removal," in Proc. Findings Assoc. Comput. Linguistics: EMNLP, 2020, pp. 2051–2061, doi: 10.18653/v1/2020.findings-emnlp.186.
[63] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in Proc. Conf. Empirical Methods Natural Lang. Process., Hong Kong, 2019, pp. 3980–3990, doi: 10.18653/v1/D19-1410.
[64] W. Wu, H. Peng, and S. Yu, "YuNet: A tiny millisecond-level face detector," Mach. Intell. Res., vol. 20, no. 5, pp. 656–665, Oct. 2023, doi: 10.1007/s11633-023-1423-y.